

Hibernate Search in Action



By : Yanai Franchi, Chief Architect , Tikal

Agenda


- ▶ The mismatch problems
- ▶ Hibernate Search in a nutshell
- ▶ Mapping – Solving the **structural** mismatch
- ▶ Indexing – Solving the **synchronization** mismatch
- ▶ Querying – Solving the **retrieval** mismatch
- ▶ Demo
- ▶ Scale Hibernate Search






The Mismatch Problems

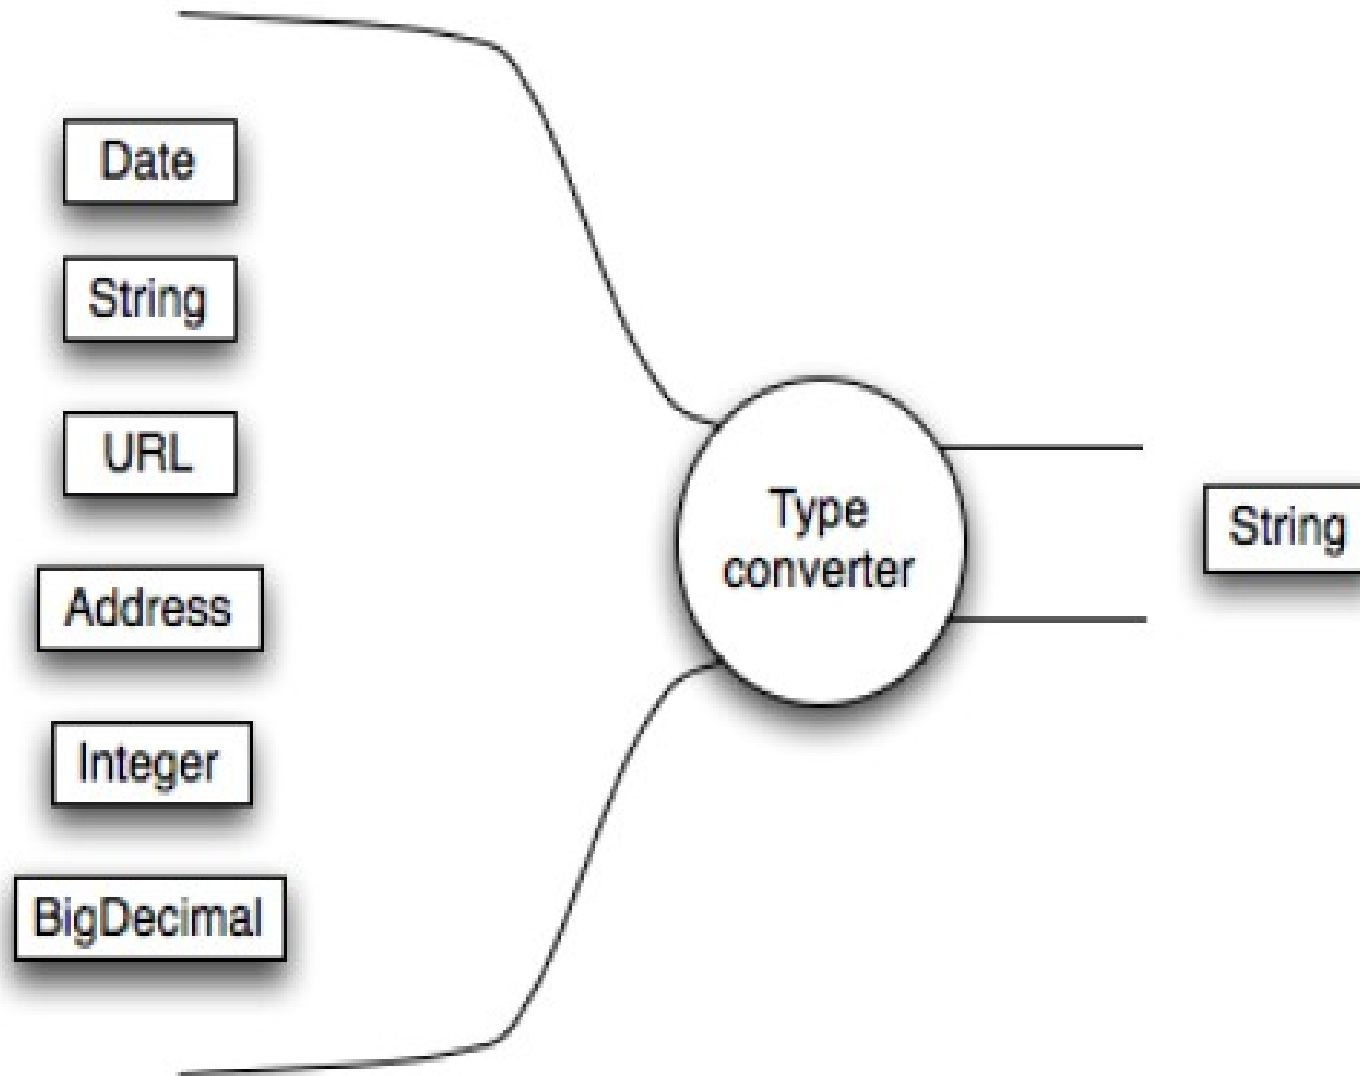
Impedance Mismatch Between Object And Index Models



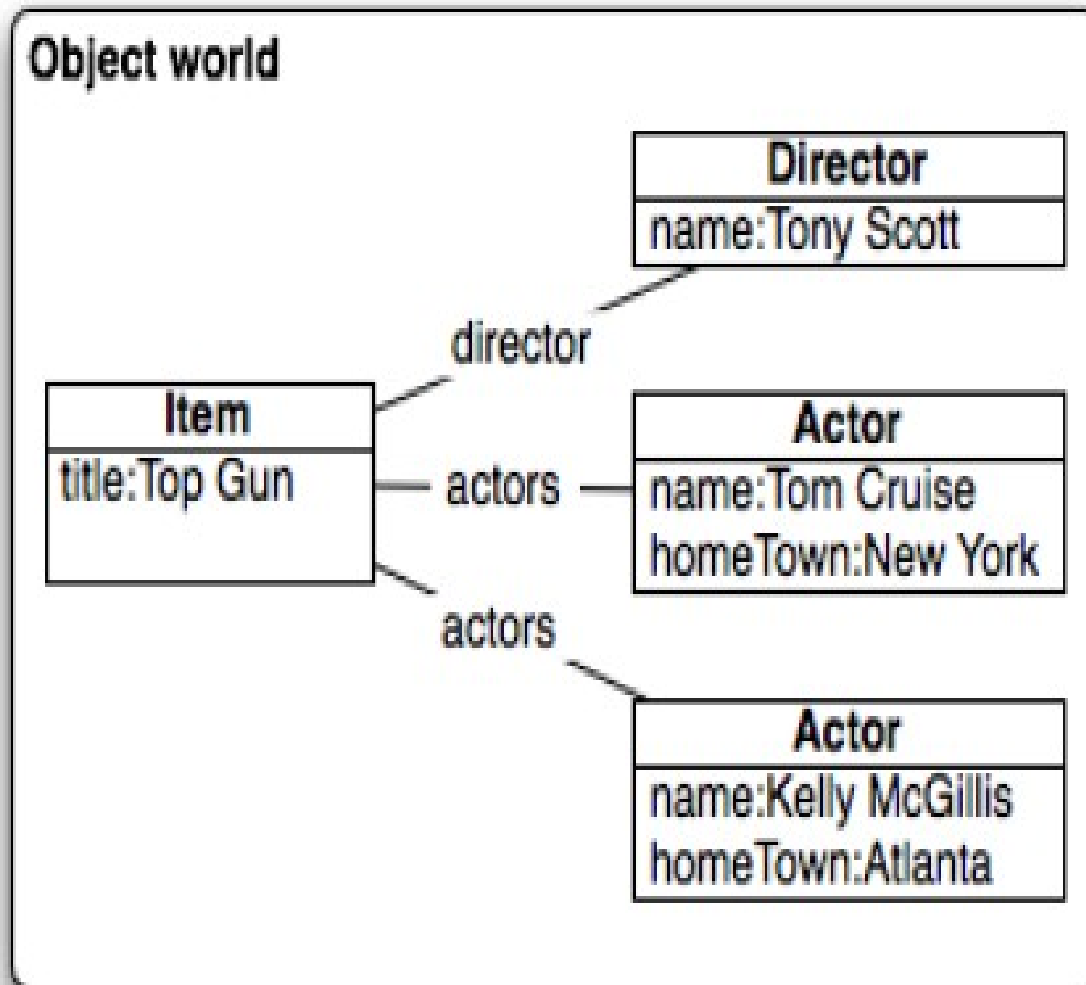
!=



Mismatch With Types



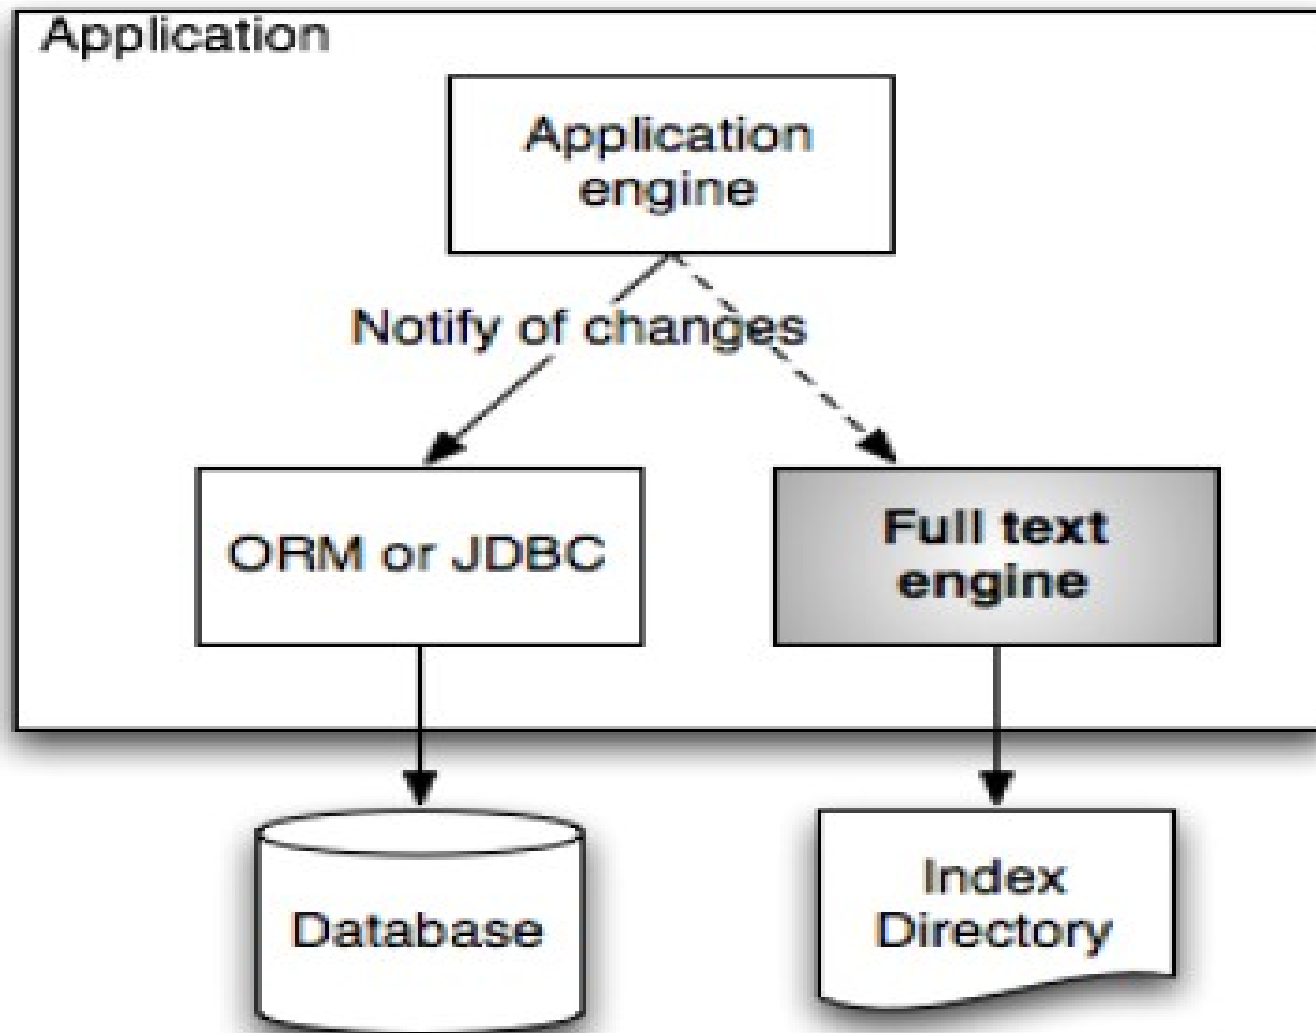
Mismatch With Associations



Index world

?

Synchronization Mismatch



Retrieval Mismatch

- ▶ NO Conversation – You don't want to go there...
 - » Loose domain driven, and OO paradigm
 - » No type safety and strong type

- ▶ Conversion
 - » “rehydrate” Document from field values stored in index.
 - No lazy loading and transparent access
 - No automatic synchronous against the DB (and index)
 - » Retrieve Hibernate managed objects.
 - Loading one-by-one is NOT efficient...





Hibernate Search in a Nutshell

Hibernate Search Goal

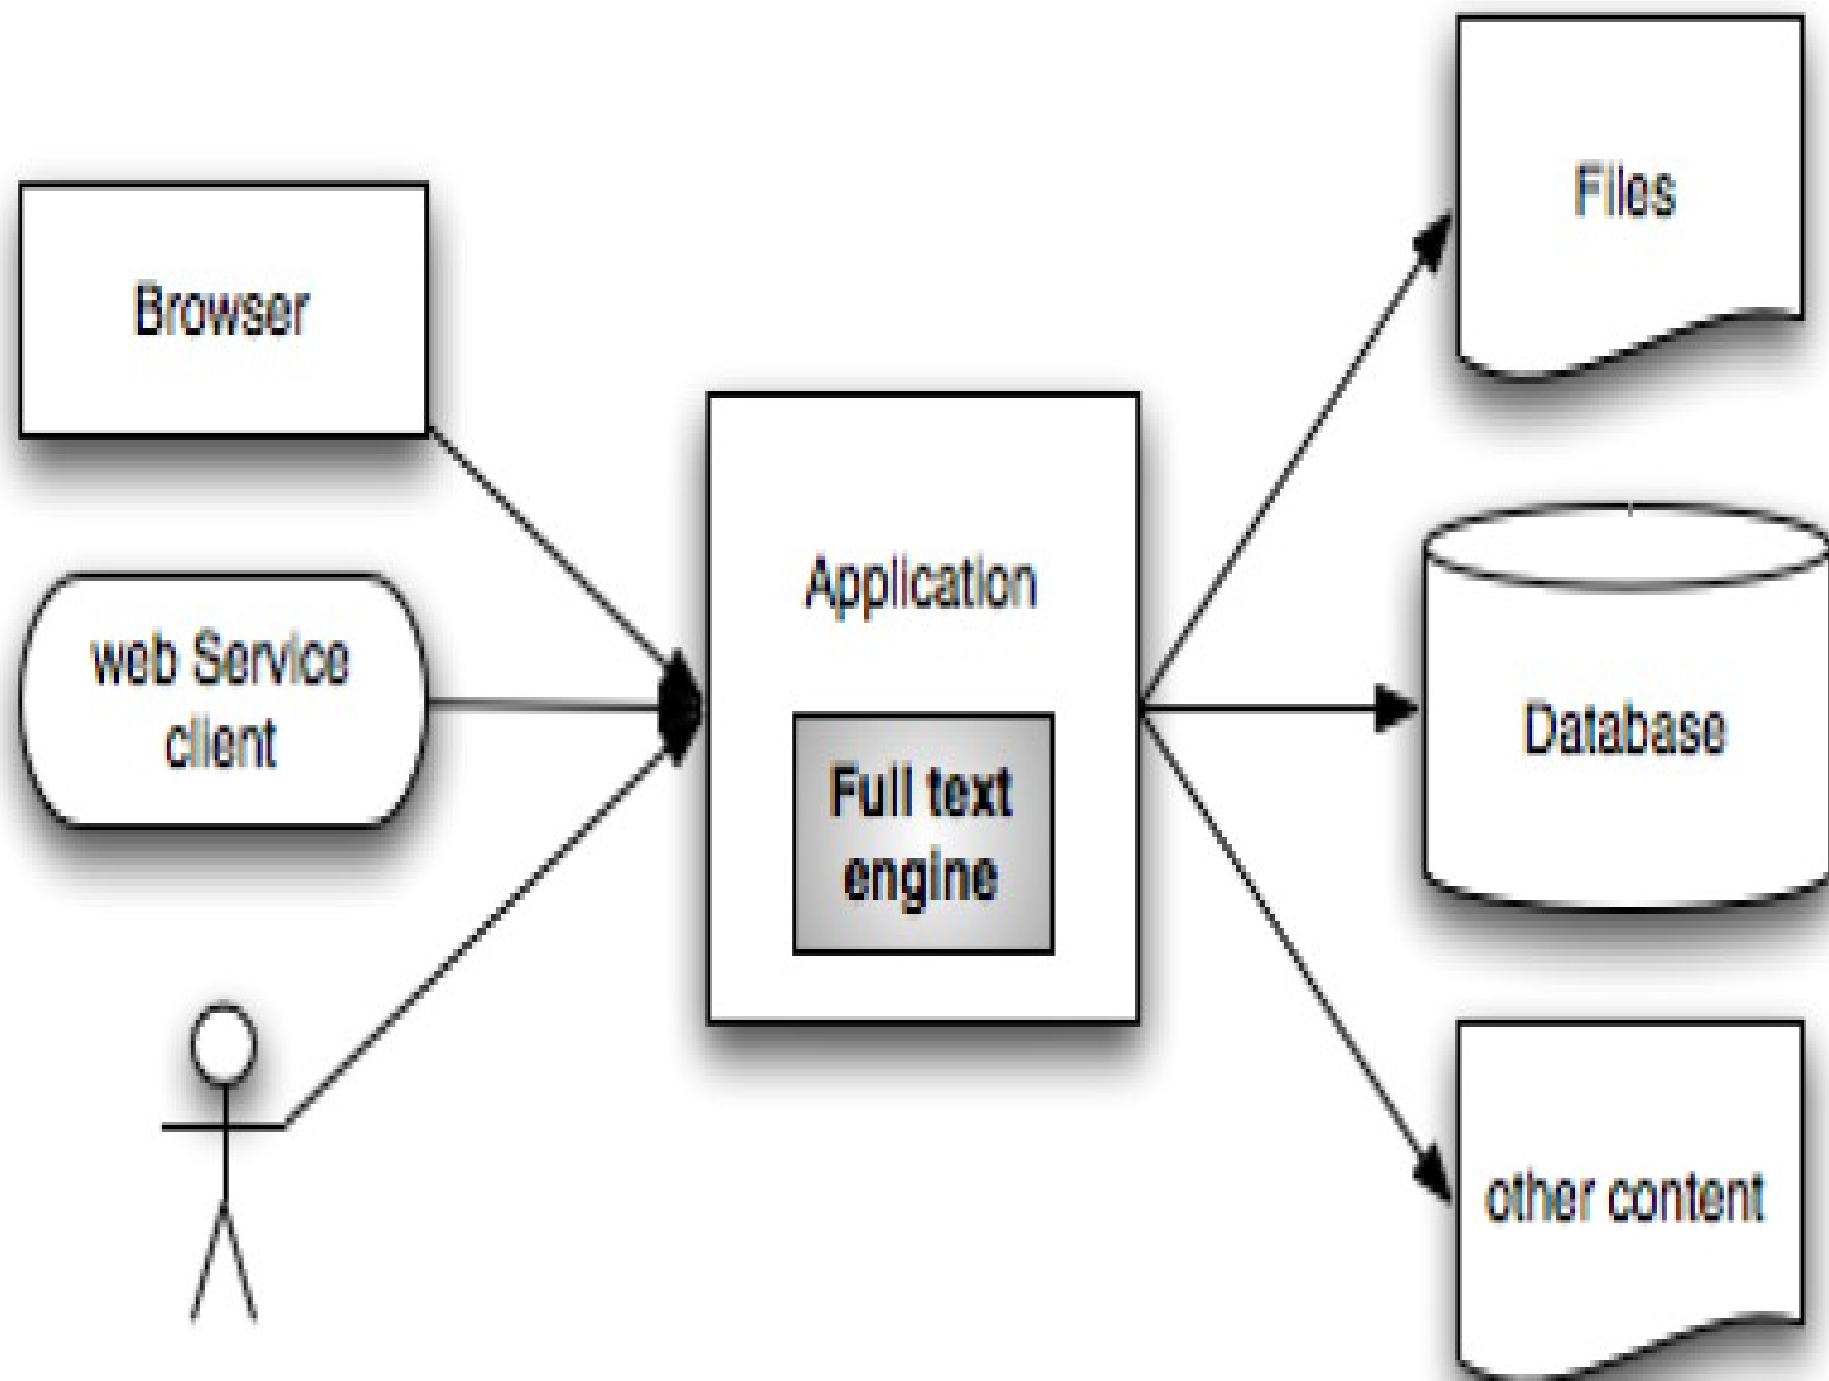
Leverage Hibernate (ORM)
and Apache Lucene (full-text search engine),
while address the mismatch problems.

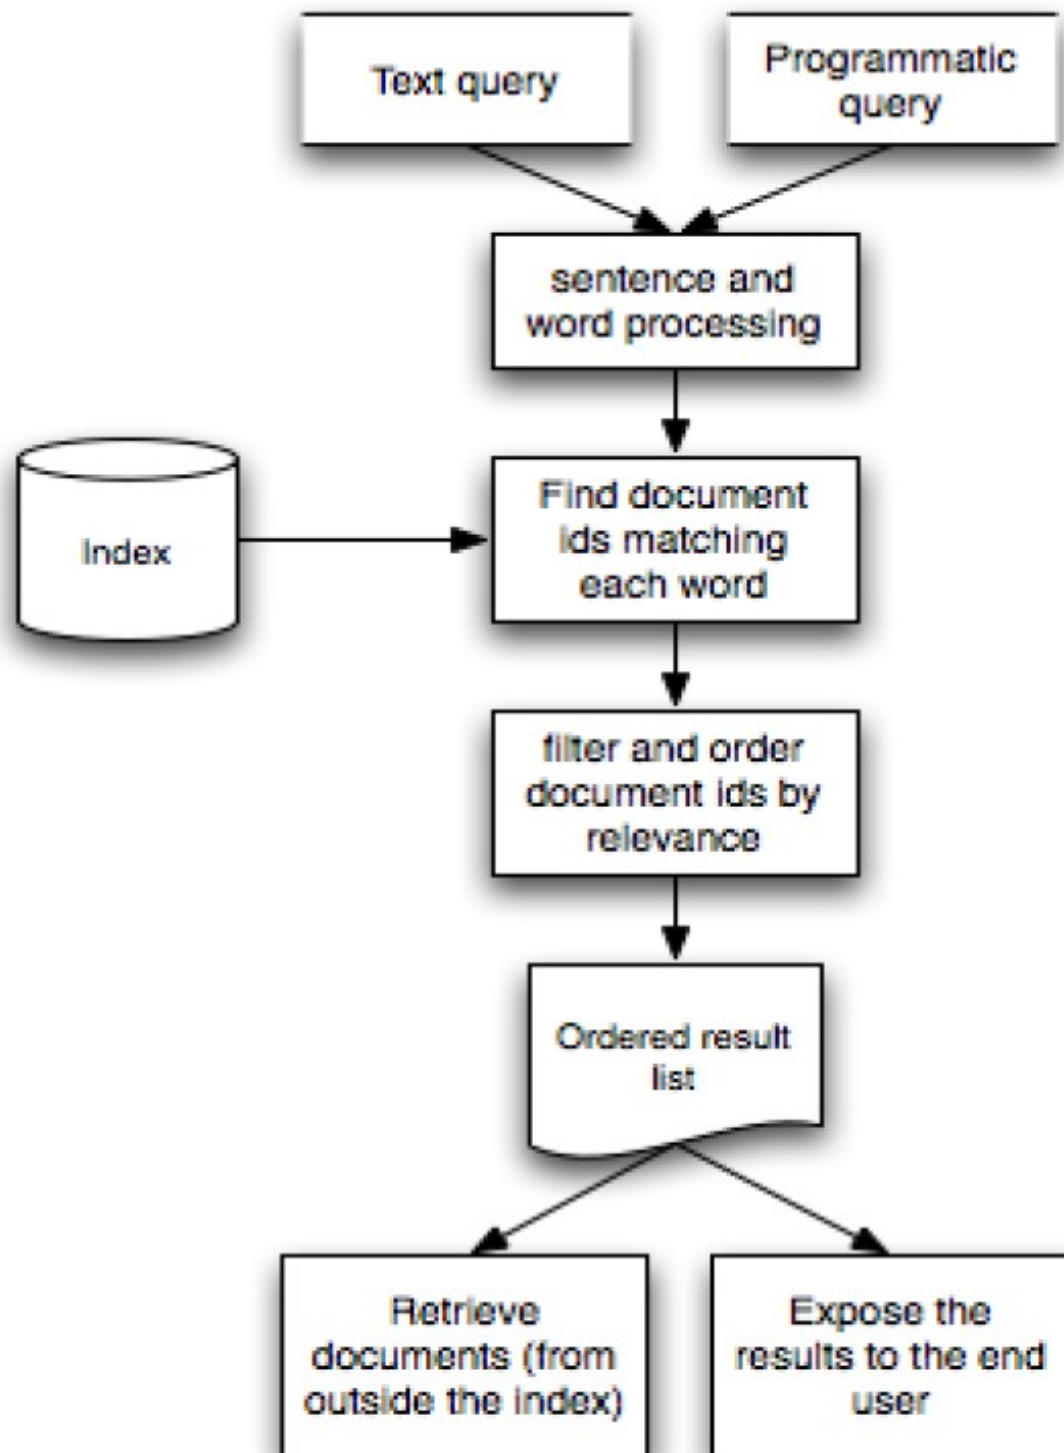


Hibernate Search Features

- ▶ Under the Hibernate platform
 - » LGPL
- ▶ Built on top of Hibernate Core
- ▶ Use Apache Lucene(tm) under the hood
 - » Hides the low level and complex Lucene API usage
- ▶ Solve the mismatches







Project Set-up

- ▶ Set your classpath
 - » hibernate-search.jar: the core API and engine of Hibernate Search
 - » lucene-core.jar: Apache Lucene engine
 - » hibernate-commons-annotations.jar: some common utilities for the Hibernate project

```
<dependency>
  <groupId>org.hibernate</groupId>
  <artifactId>hibernate-search</artifactId>
  <version>3.1.0.GA</version>
</dependency>
<dependency>
  <groupId>org.hibernate</groupId>
  <artifactId>hibernate-annotations</artifactId>
  <version>3.4.0.GA</version>
</dependency>
```

(Optional) Project Configuration

- ▶ Configure hibernate search
 - » No need for event listeners.
 - When using JPA/Hibernate Annotations
- ▶ hibernate-cfg.xml or META-INF/persistence.xml

```
<?xml version="1.0" encoding="UTF-8"?> META-INF/persistence.xml
<persistence>
  <persistence-unit name="dvdstore-catalog">
    <jta-data-source>jdbc/test</jta-data-source>
    <properties>
      ...
      <property
        name="hibernate.search.default.indexBase"
        value="/users/application/indexes"/>
      ..
    </properties>
  </persistence-unit>
</persistence>
```


Map Your Domain Model

```
@Entity
@Indexed
public class Book {
    @Id // → Automatically mapped to @DocumentId
    private Integer id;

    @Field
    private String title;

    ...
}
```

How Is The Index Look Like?

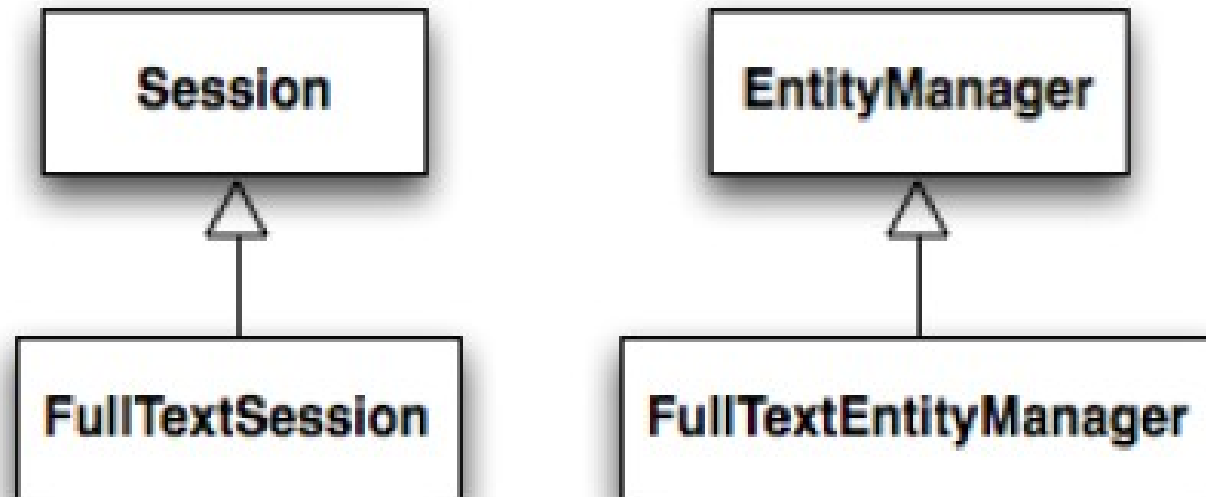
Book table

id	title
1	The Da Vinci Code
2	Alice's adventures in Wonderland
...	...
234	Economics for Dummies
235	The real history behind the Da Vinci Code

Full text index

Word	id	freq.	position
adventures	2	1/3	2
alice	2	1/3	1
code	1	1/3	3
	235	1/6	5
da	1	1/3	1
	235	1/6	4
dummies	234	1/2	2
economics	234	1/2	1
vinci	1	1/3	2
	235	1/6	5
wonderland	2	1/3	3

Hibernate Search Managers



```
Session session = sessionFactory.getCurrentSession();

FullTextSession fts =
    org.hibernate.search.Search.getFullTextSession(session);
```

```
@PersistenceContext EntityManager em;
...

FullTextEntityManager ftem =
    org.hibernate.search.jpa.Search.getFullTextEntityManager(em);
```

Query in Action

```
String searchStr = "title:hypernate~ OR description:persistence";

org.apache.lucene.search.Query luceneQuery =
    buildLuceneQuery(searchStr);

javax.persistence.Query jpaQuery =
    ftEm.createFullTextQuery(luceneQuery, Book.class, Course.class);

List booksAndCourses = query.getResultList();

applySomeChanges(booksAndCourses);
```

Can accept
more than
one class


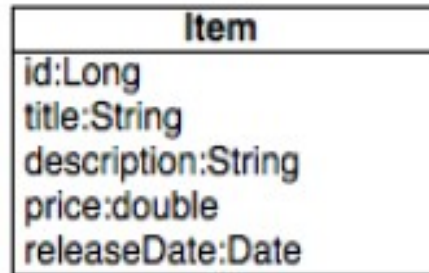
Books and Courses get into JPA “*persistent context*”
and changes will be automatically applied to **DB** (and
the *Lucene Index*)



Mapping – Solve The Structural Mismatch



Object world



Additional types


Long

String

double

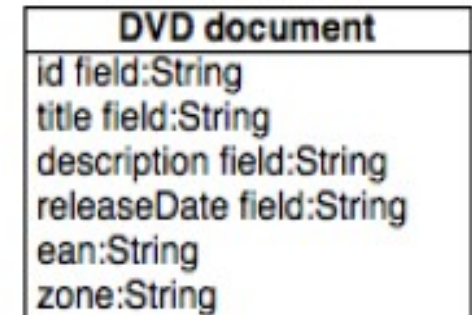

Date

Zone <<Enum>>



Conversion to
(and
sometimes
from)

Lucene index world



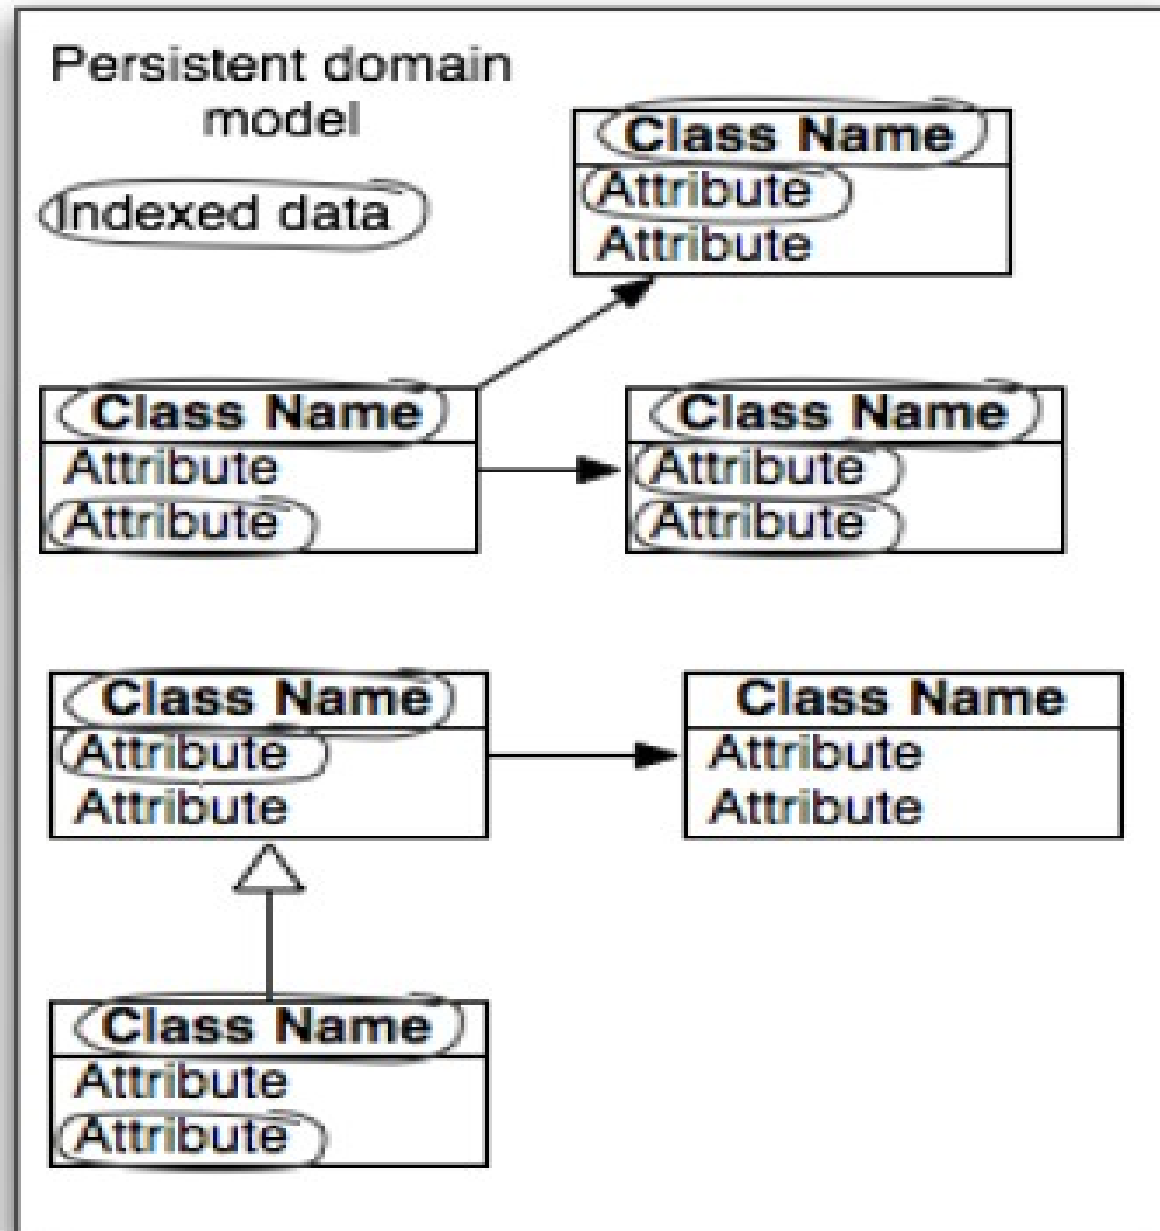
Additional types

String

Mapping Entity & Primary Key

```
@Entity
@Indexed
public class Item {
    @Id // → Automatically mapped to @DocumentId
    @GeneratedValue
    private Integer id;
    ...
}
```


Marking Properties As Indexed



Mapping Properties

```
@Entity
@Indexed
public class Item {
    @Id @GeneratedValue
    private Integer id;
    @Field(index=Index.UN_TOKENIZED)
    private String ean;

    @Field(store=Store.YES)
    private String title;

    //Will not be indexed while still being stored into DB
    private String imageURL;

    private String description;
    ...
    @Field //Annotation on the getter
    public String getDescription() {
        return this.description;
    }
}
```

Multiple Indexed Property

```
@Entity
@Indexed
public class Item {
    ...
    @Fields({
        @Field(index=Index.TOKENIZED)
        @Field(name="title_sort", index=Index.UN_TOKENIZED)
    })
    private String title;
```

- ▶ Properties that will be used to sort query results (rather than by relevance) must not be tokenized but must be indexed.
 - » Use UN_TOKENIZED indexing strategy

Mapping Inheritance

```
@Entity //Superclasses do not have to be marked @Indexed
public abstract class Item {
    @Id // used as @DocumentId
    @GeneratedValue
    private Integer id

    @Field //Superclasses can contain indexed properties
    private String title;
    ...
}
```

```
@Entity
@Indexed //Concrete subclasses are marked @Indexed
public class Dvd extends Item {

    @Field(index=Index.UN_TOKENIZED)
    private String ean;
    ...
}
```

Built-In Bridges

- ▶ Bridges convert a Java object type into a string.
- ▶ Some field bridges also convert back the string into the original object structure
 - » Identity and projected fields
- ▶ Hibernate Search comes with many out-of-the-box field bridges. But you can write (or reuse) you own...
 - » PDF, Microsoft-Word and other document types
 - » Index Year, Month, Day on separate fields
 - » Make numbers comparable



Bridge Issues

► Dates

- » [20080112 TO 20080201] – field is between 12 January 2008 and 1 February 2008.
- » Hibernate Search lets you pick the date precision you wish from year to millisecond:

```
@DateBridge( resolution = Resolution.DAY )  
private Date birthdate;  
  
@DateBridge( resolution = Resolution.MINUTE )  
private Date flightArrival;
```

► Numbers

- » “2 > “12”
- » [6 TO 9] => 6 OR 7 OR 8 OR 9

Custom Bridge in Action

```
@Entity @Indexed
public class Item {
    @Field
    @FieldBridge(
        impl=PaddedRoundedPriceBridge.class, // So 2 becomes "002"
        params= { @Parameter(name="pad", value="3") }
    )
    private double price;
    ...
}
```

- Mapping a property to split the information to multiple fields in the index.

ClassBridge

```
@Entity @Indexed
@ClassBridge(name="childrenOnly",
              impl=ChildrenFlagBridge.class, index=Index.UN_TOKENIZED)
public class Item {
    ...
}
```

```
public class ChildrenFlagBridge implements StringBridge {
    public String objectToString(Object object) {
        Item item = (Item) object;
        Category childrenCategory = new Category("Children");

        boolean hasChildrenCategory =
            item.getCategories().contains(childrenCategory);

        return hasChildrenCategory ? "yes" : "no";
    }
}
```


How to Map Associations ?

Entity

Relation


Entity

Query






1


*




Navigation
and filter



SQL join
and where clause

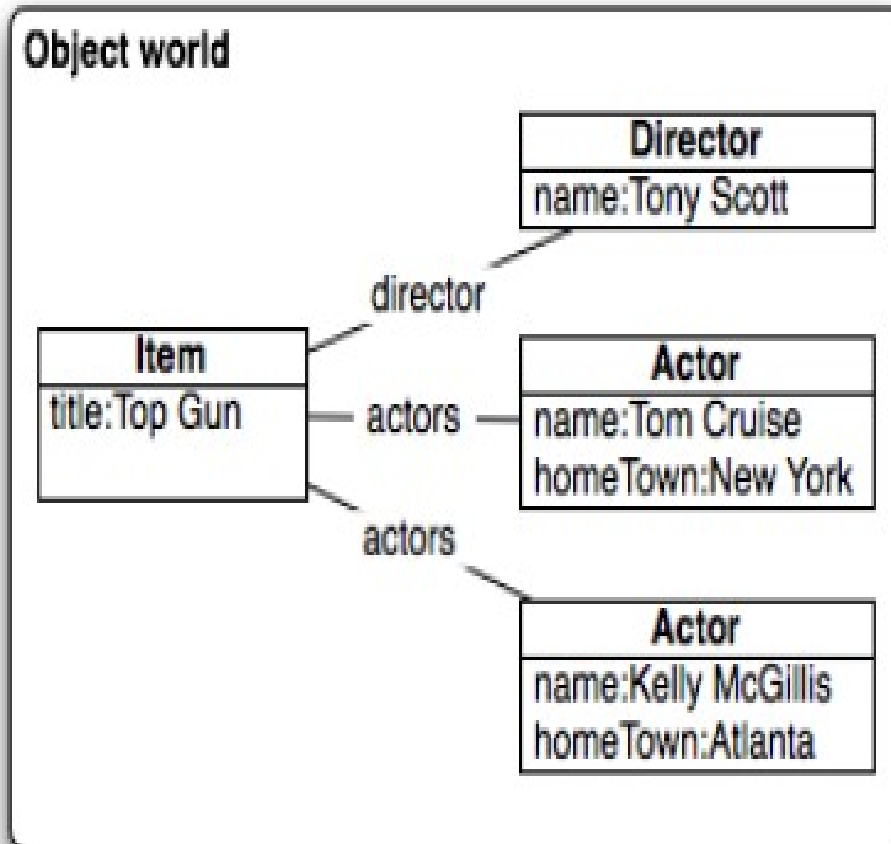


?



?

De-Normalize Associations



Lucene index world

Item document
title:Top Gun
director.name:Tony Scott
actors.name:Tom Cruise Kelly McGillis
actors.homeTown: New York Atlanta

De-Normalization Implication

- ▶ Can return items that :
 - » One of the actor is “**Cruise**” and another one is “**McGillis**”
 - » One of the actor is either “**Cruise**” or “**McGillis**”
 - » “**Cruise**” plays but not “**McGillis**”
- ▶ Can ****NOT**** do:
 - » Return items where one of the actor is “**Tom**” and his home town is “**Atlanta**”.
 - Turn the query upside down by targeting actor as the root entity and then collect the matching items
 - Use a query filter to refine an initial query
- ▶ Sometime you may end up in a dead end...
 - » Apply part of the query (the discriminant part) in Lucene,
 - » Collect the matching identifiers
 - » Run a HQL query restricting by these identifiers.



Indexing Embeddables

```
@Embeddable
public class Rating {
    @Field(index=Index.UN_TOKENIZED) private Integer overall;
    @Field(index=Index.UN_TOKENIZED) private Integer scenario;
    @Field(index=Index.UN_TOKENIZED) private Integer soundtrack;
    @Field(index=Index.UN_TOKENIZED) private Integer picture;
    ...
}
```

```
@Entity @Indexed
public class Item {
    @IndexedEmbedded private Rating rating;
    ...
}
```

“find items with overall rating equals to 9”

rating.overall : 9

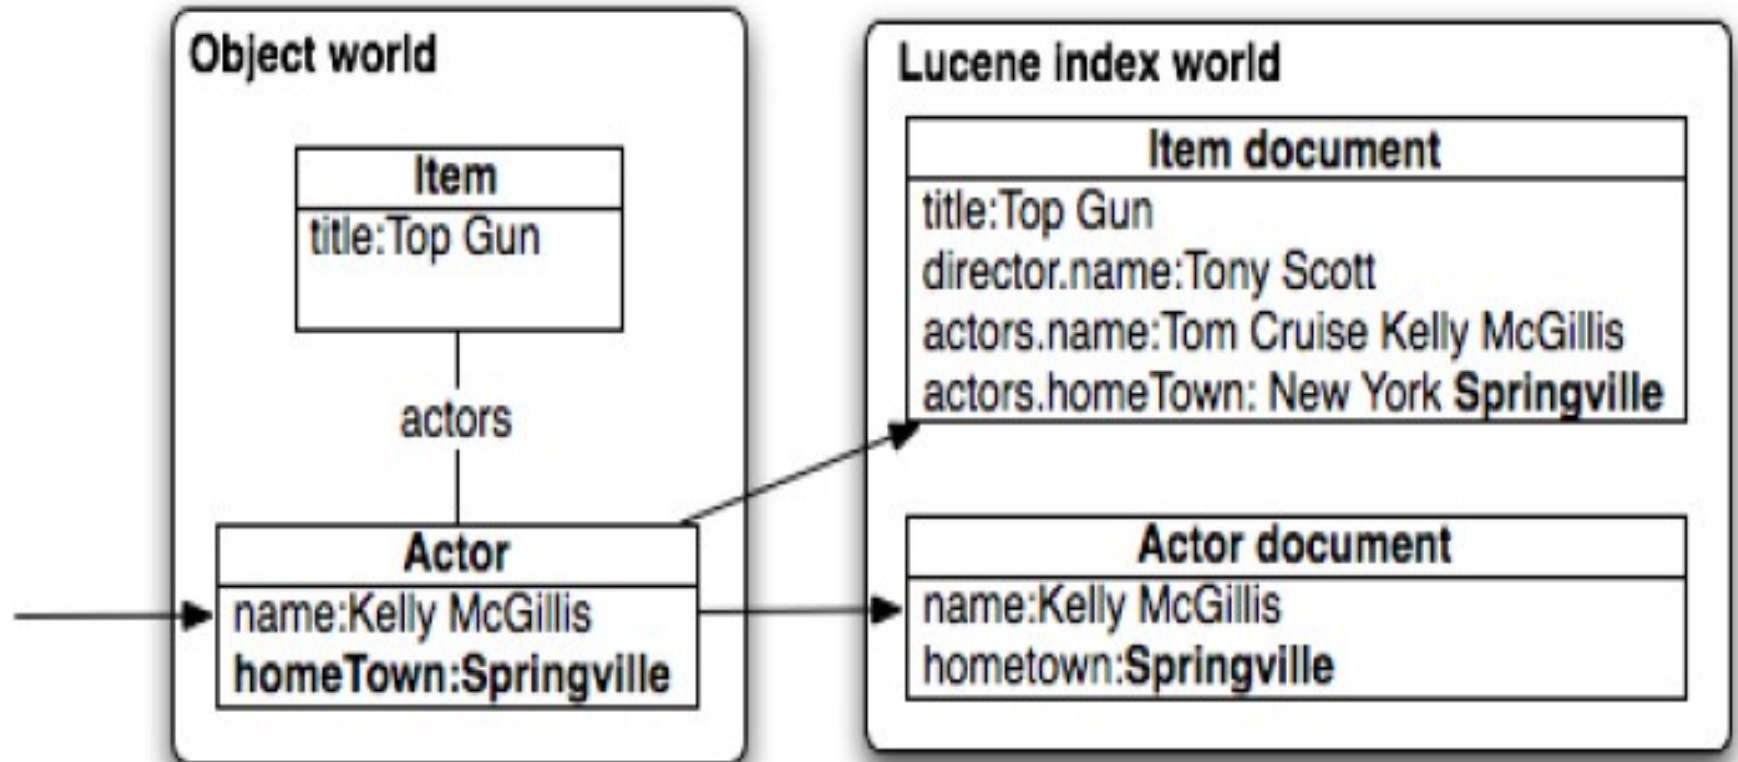
...And Embeddables Collection

```
@Embeddable
public class Country {
    @Field private String name;
    ...
}
```

```
@Entity @Indexed
public class Item {
    @CollectionOfElements @IndexedEmbedded
    private Collection<Country> distributedIn;
    ...
}
```

- Don't abuse IndexedEmbedded. Be careful on collection indexing...

Indexing Associated Entities



→ Propagated change

When a change is done on an associated entity, Hibernate Search must update all the documents where the entity is embedded in

Indexing Associated Entities

```
@Entity @Indexed
public class Item {
    @ManyToMany
    @IndexedEmbedded
    private Set<Actor> actors; //embed actors when indexing
    ...
}
```

```
@Entity @Indexed
public class Actor {
    @Field private String name;

    @ManyToMany(mappedBy="actors")
    @ContainedIn // We may use (depth=4) to limit depth
    private Set<Item> items;
    ...
}
```

- Relations between entities become bi-directional in case the Actor is not immutable, or do manual index



Indexing Your Data - Solve The Synchronization Mismatch

Defining a DirectoryProvider

```
# Production configuration
hibernate.search.default.directory_provider
    org.hibernate.search.store.FSDirectoryProvider
hibernate.search.default.indexBase /User/production/indexes
```

```
# File directory structure
/Users
  /Production
    /indexes
      /com.manning.hsia.dvdstore.model.Item
      /com.manning.hsia.dvdstore.model.Actor
```

```
# Test Configuration
hibernate.search.default.directory_provider
    org.hibernate.search.store.RAMDirectoryProvider
```

Analizers – Lucene Brain

- ▶ The key feature of the full text search
- ▶ Taking text as an **input**, **chunking** it into individual words (by a tokenizer) and optionally applying some **operations (by filters)** on the tokens.
- ▶ Applied: globally, per entity, or per property



Tokenizers & Filters

- ▶ StandardTokenizer – **Splits words** at punctuation characters and **removing punctuation** signs with a couple of exception rules.
- ▶ Filters alter the stem of tokens (remove/change/add)
 - » **StandardFilter** – Removes apostrophes and acronyms dots
 - » **LowerCaseFilter**
 - » **StopFilter** – Eliminates “noise” words.




StandardAnalyzer in Action

```
@AnalyzerDef( //This is the default → no need to write it
    name="applicationAnalyzer",
    tokenizer =@TokenizerDef(factory=StandardTokenizerFactory.class),
    filters = {
        @TokenFilterDef(factory = StandardFilterFactory.class),
        @TokenFilterDef(factory = LowerCaseFilterFactory.class),
        @TokenFilterDef(factory = StopFilterFactory.class)
    }
)
```

```
@Entity @Indexed
@Analyzer(definition="applicatioAanalyzer")
public class Item {
    ...
}
```

More Available Filters

Synonym



Stem

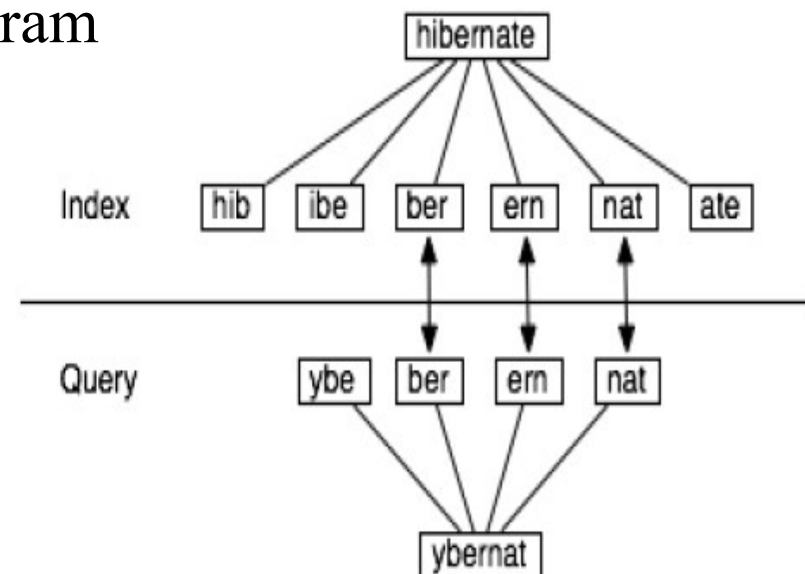
Rule			Example		
SSSES	->	SS	Caresses	->	Caress
IES	->	I	Ponies	->	Poni
SS	->	SS	Caress	->	Caress
S	->		Cats	->	Cat

Phonetic



Les Misérable => LS MSRP

N-Gram



N-Gram Analyzer Example

```
@AnalyzerDef(  
    name="ngramAnalyzer",  
    tokenizer =@TokenizerDef(factory=StandardTokenizerFactory.class),  
    filters = {  
        //Standard, LowerCase and Stop filters goes here  
        @TokenFilterDef(factory = NGramTokenFilterFactory.class,  
            params = {  
                @Parameter(name="minGramSize", value="3"),  
                @Parameter(name="maxGramSize.", value="3")  
            })  
    }  
)
```

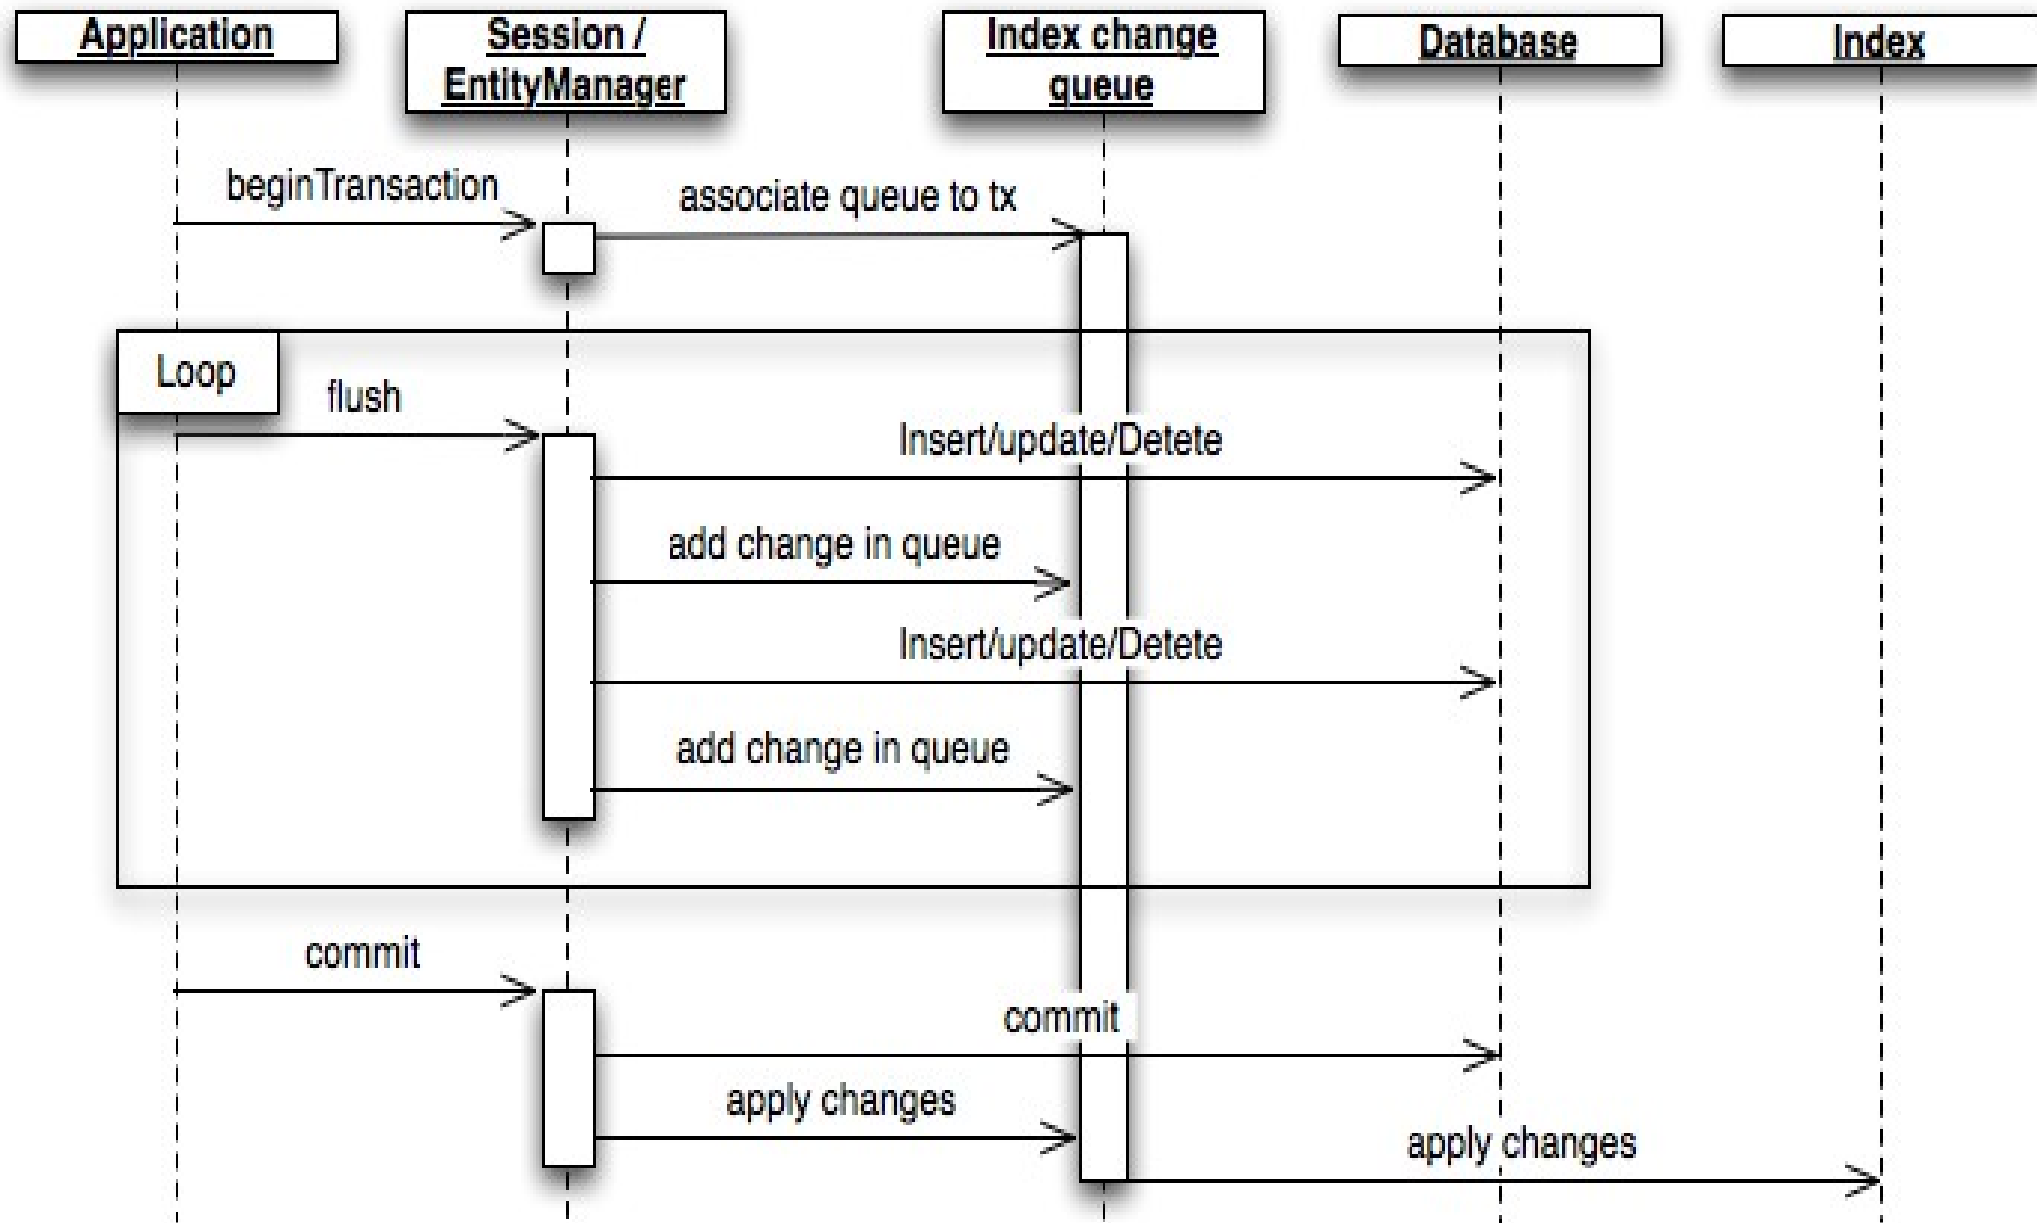
```
@Entity @Indexed  
// The default StandardAnalyzer will be used  
public class Item{  
    @Fields({  
        @Field(index=Index.TOKENIZED),  
        @Field(name="title_ngram",index=Index.UN_TOKENIZED,  
            analyzer=@Analyzer(definition="ngramAnalyzer")  
    })  
    private String title;
```

Which Technique to Choose?

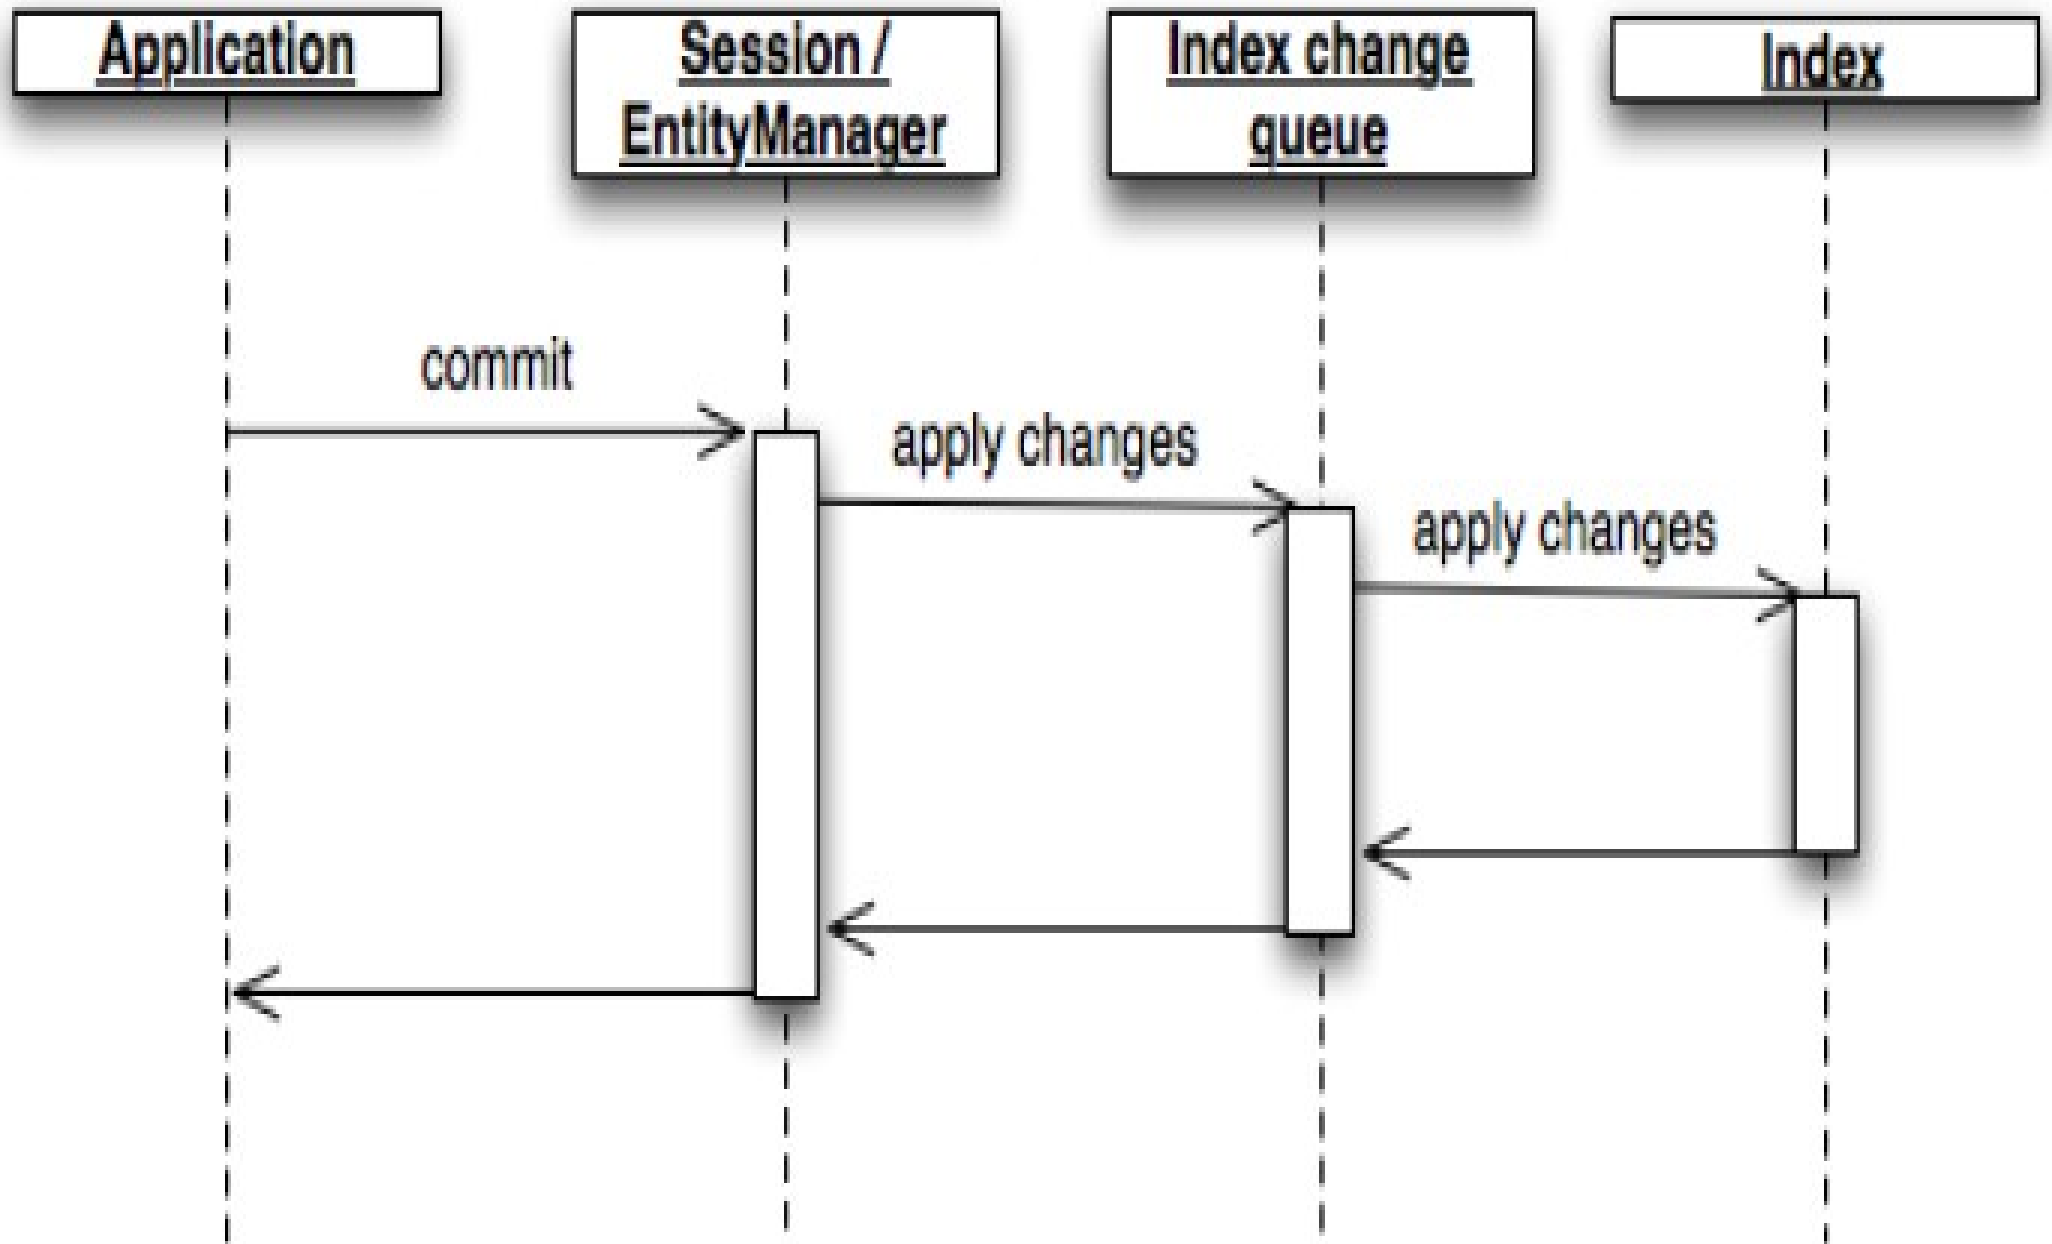
- ▶ Use approximation analyzers on dedicated fields.
- ▶ Search in layers – Expand the approximation level.
 - » The search engine can execute the strict query first
 - » If more data is required a second query using approximation techniques can be used and so on.
 - » Once the search engine has retrieved enough information, it bypasses the next layers.
- ▶ Remember that a Lucene query is quite cheap. Running several Lucene queries per user query is perfectly acceptable.



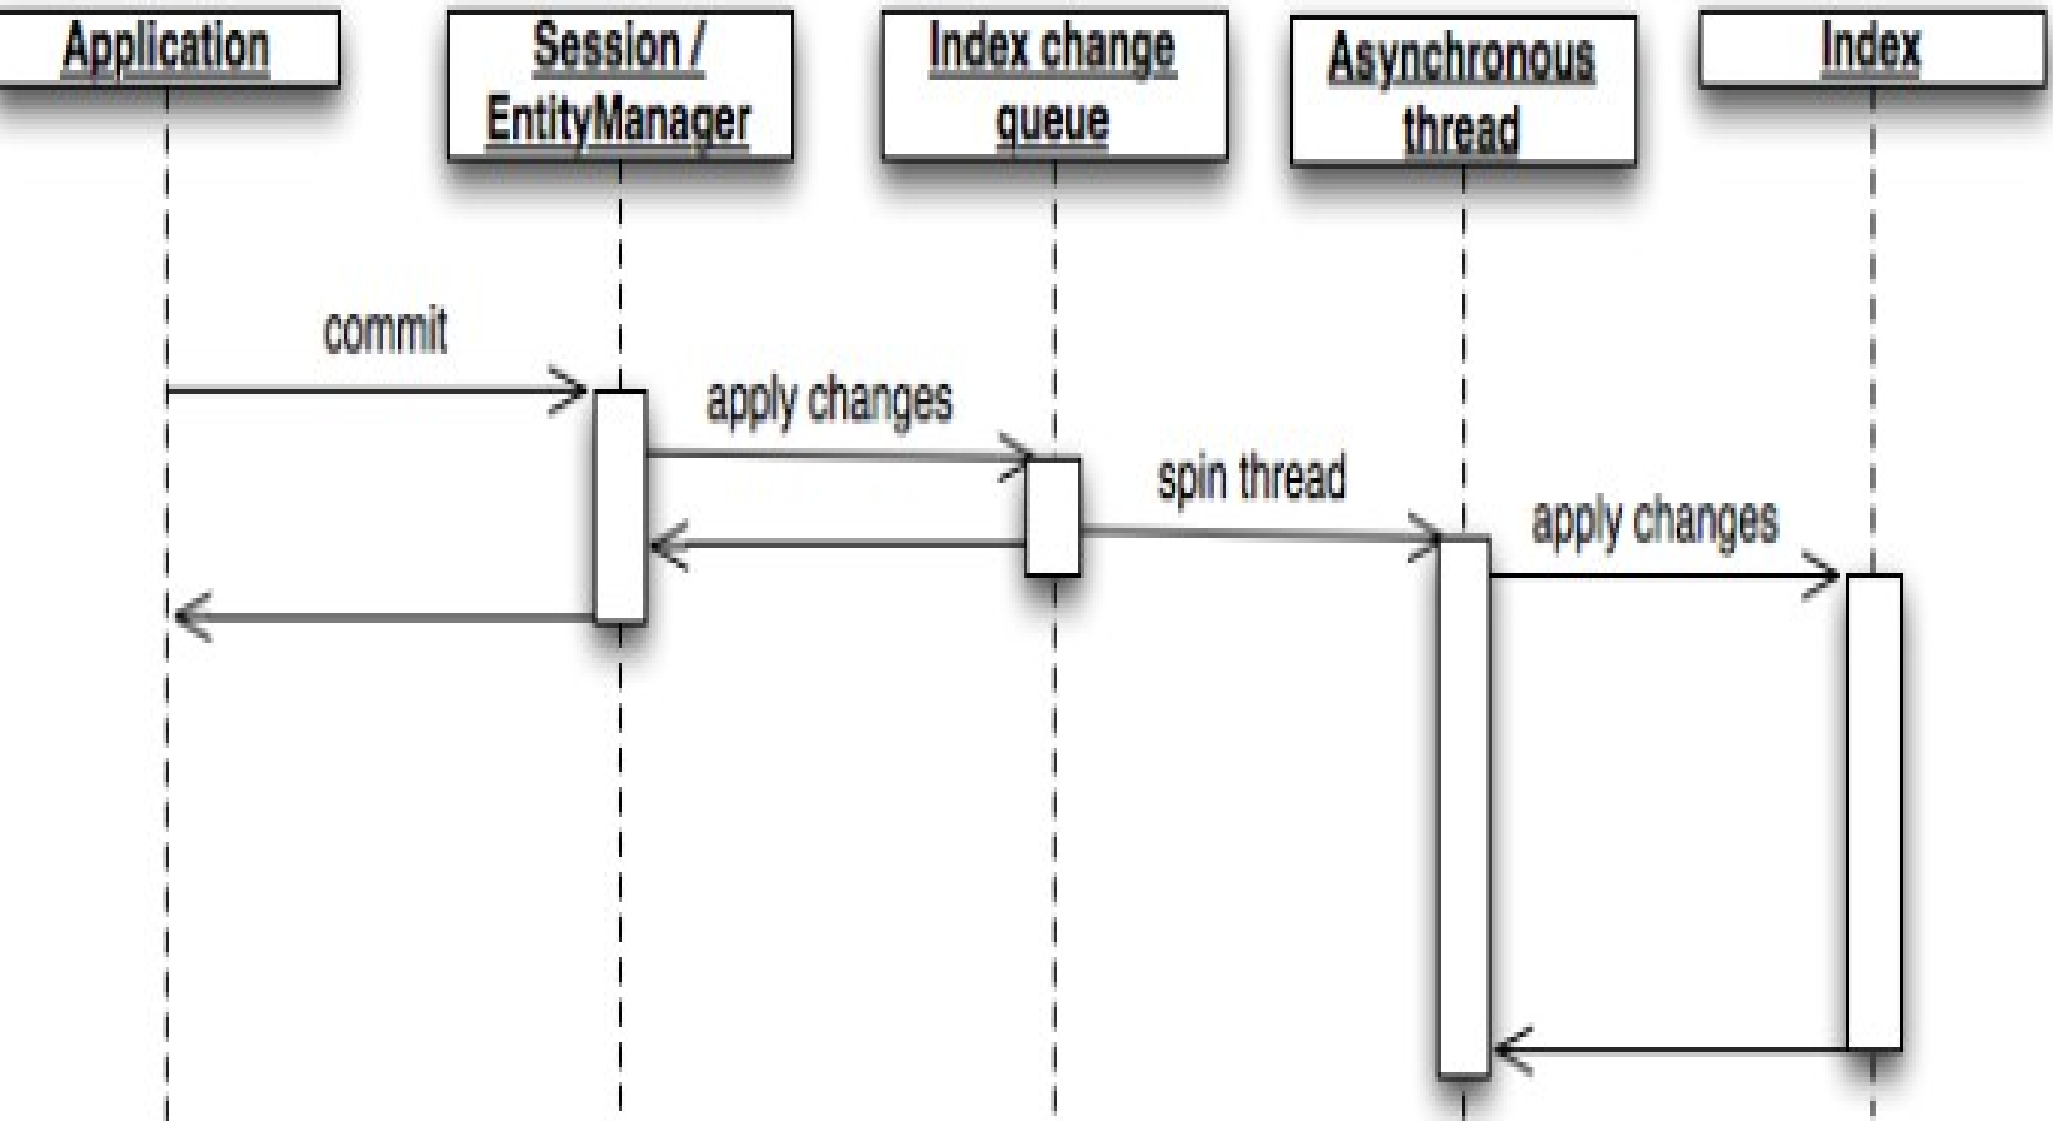
Indexing Flow Diagram



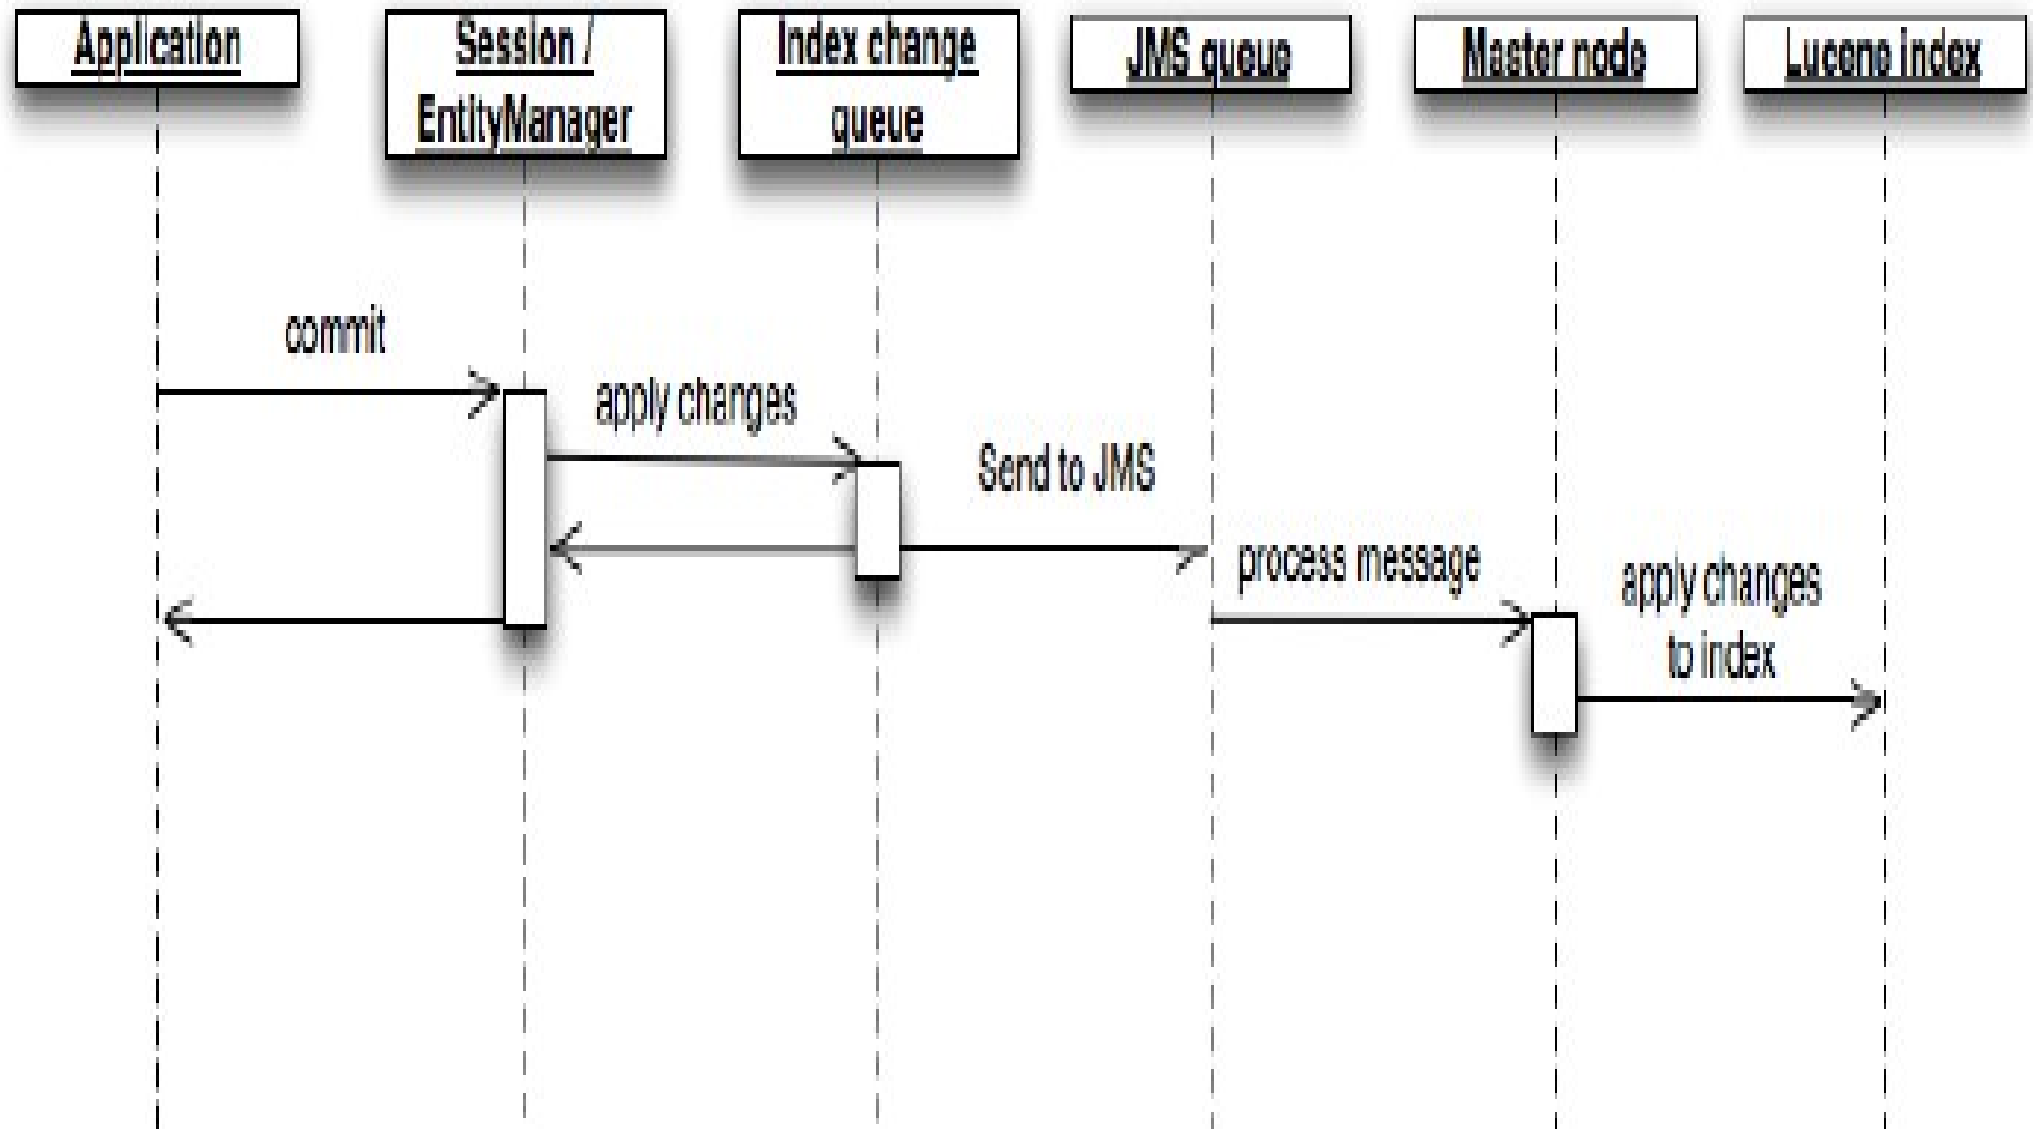
Synchronous Flow



Asynchronous Flow



JMS Flow



Manual Index – Naïve Approach

```
Transaction tx = session.beginTransaction();

//read the data from the database
Query query = ftSession.createCriteria(Item.class);
List<Item> items = query.list();

//index the data
for (Item item : items) {
    ftSession.index(item);
}

tx.commit();
```

OutOfMemoryError

Load “distributor”
for each item

Manual Index - The Right Way


```
Transaction tx = ftSession.beginTransaction();

ftSession.setFlushMode(FlushMode.MANUAL); //disable flush
ftSession.setCacheMode(CacheMode.IGNORE); //disable 2nd level cache

ScrollableResults results = ftSession.createCriteria( Item.class )
    .setFetchMode("distributor", FetchMode.JOIN)
    .setResultTransformer(CriteriaSpecification.DISTINCT_ROOT_ENTITY)
    .setFetchSize(BATCH_SIZE);
    .scroll( ScrollMode.FORWARD_ONLY );


for(int i=1; results.next() ; i++) {
    ftSession.index( results.get(0) );
    if (i % BATCH_SIZE == 0) {
        ftSession.flushToIndexes(); //apply changes to the index
        ftSession.clear(); //clear the session releasing memory
    }
}
tx.commit(); //apply the remaining index changes
```

Index With Batch Approach




`hibernate.search.indexing_strategy = manual`

Mix Batch And Event Approach



Third Party Updates Your DB



What Influences Indexing Time

- ▶ Number of properties indexed
- ▶ Type of analyzer used
- ▶ Properties stored
- ▶ Properties embedded
- ▶ On Mass Indexing
 - » Index asynchronously
 - » Index on a different machine
 - » Use our previous manual sample as a template
 - » `session.getSearchFactory().optimize();`





Query – Solving The Retrieval Mismatch

Full-Text Search Query

- ▶ Running Hibernate-Search Query:
 - » Building a Lucene query to express the full text search (either through the query parser or the programmatic API)
 - » Building an Hibernate Search query wrapping the Lucene query
 - » Execute Hibernate Search Query.

- ▶ But why do we need this wrapper around Lucene ?
 - » Build the Lucene Query is easy :
 - title:Always description:some desc actors.name:Tom Cruise

Executing Lucene Query Is Low Level API

- ▶ Open the Lucene directory(ies)
- ▶ Build one or several IndexReaders, and an IndexSearcher on top of them
- ▶ Call the appropriate execution method from IndexSearcher.
- ▶ Resource management for Lucene API
- ▶ Convert Documents into objects of your domain model.
 - » “rehydrate” values from Lucene index
 - No lazy loading, No transparent access, No change propagation
 - » Load entities using ORM
 - Loading one by one will work inefficiently

Hibernate Search Query

- ▶ Return managed Hibernate entities.
- ▶ Query API is similar. Use the same Query API as JPA or Hibernate-Query API.
- ▶ Query semantic is also similar.
 - » Lazy loading mechanism.
 - » Transparent propagation to DB and Index



Build Lucene Query With QueryParser

```
private org.apache.lucene.search.Query buildLuceneQuery
    (String words, Class<?> searchedEntity) {

    Analyzer analyzer = getFTEntityManager().getSearchFactory()
        .getAnalyzer(searchedEntity);

    QueryParser parser = new QueryParser( "title", analyzer );
    org.apache.lucene.search.Query luceneQuery = parser.parse(words);
    return luceneQuery;
}
```

Build Lucene Query With MutilFieldParser

```
private org.apache.lucene.search.Query buildLuceneQuery
    (String words, Class<?> searchedEntity) {

    Analyzer analyzer=getFTEntityManager().getSearchFactory()
        .getAnalyzer(searchedEntity);

    String[] productFields = {"title", "description"};
    Map<String,Float> boostPerField = new HashMap<String,Float>;
    boostPerField.put( "title", 4f);
    boostPerField.put( "description", 1f);

    QueryParser parser = new MultiFieldQueryParser(
        productFields,analyzer,boostPerField);

    org.apache.lucene.search.Query luceneQuery = parser.parse(words);
    return luceneQuery;
}
```

Build & Execute The FullTextQuery

```
@PersistenceContext private EntityManager em;

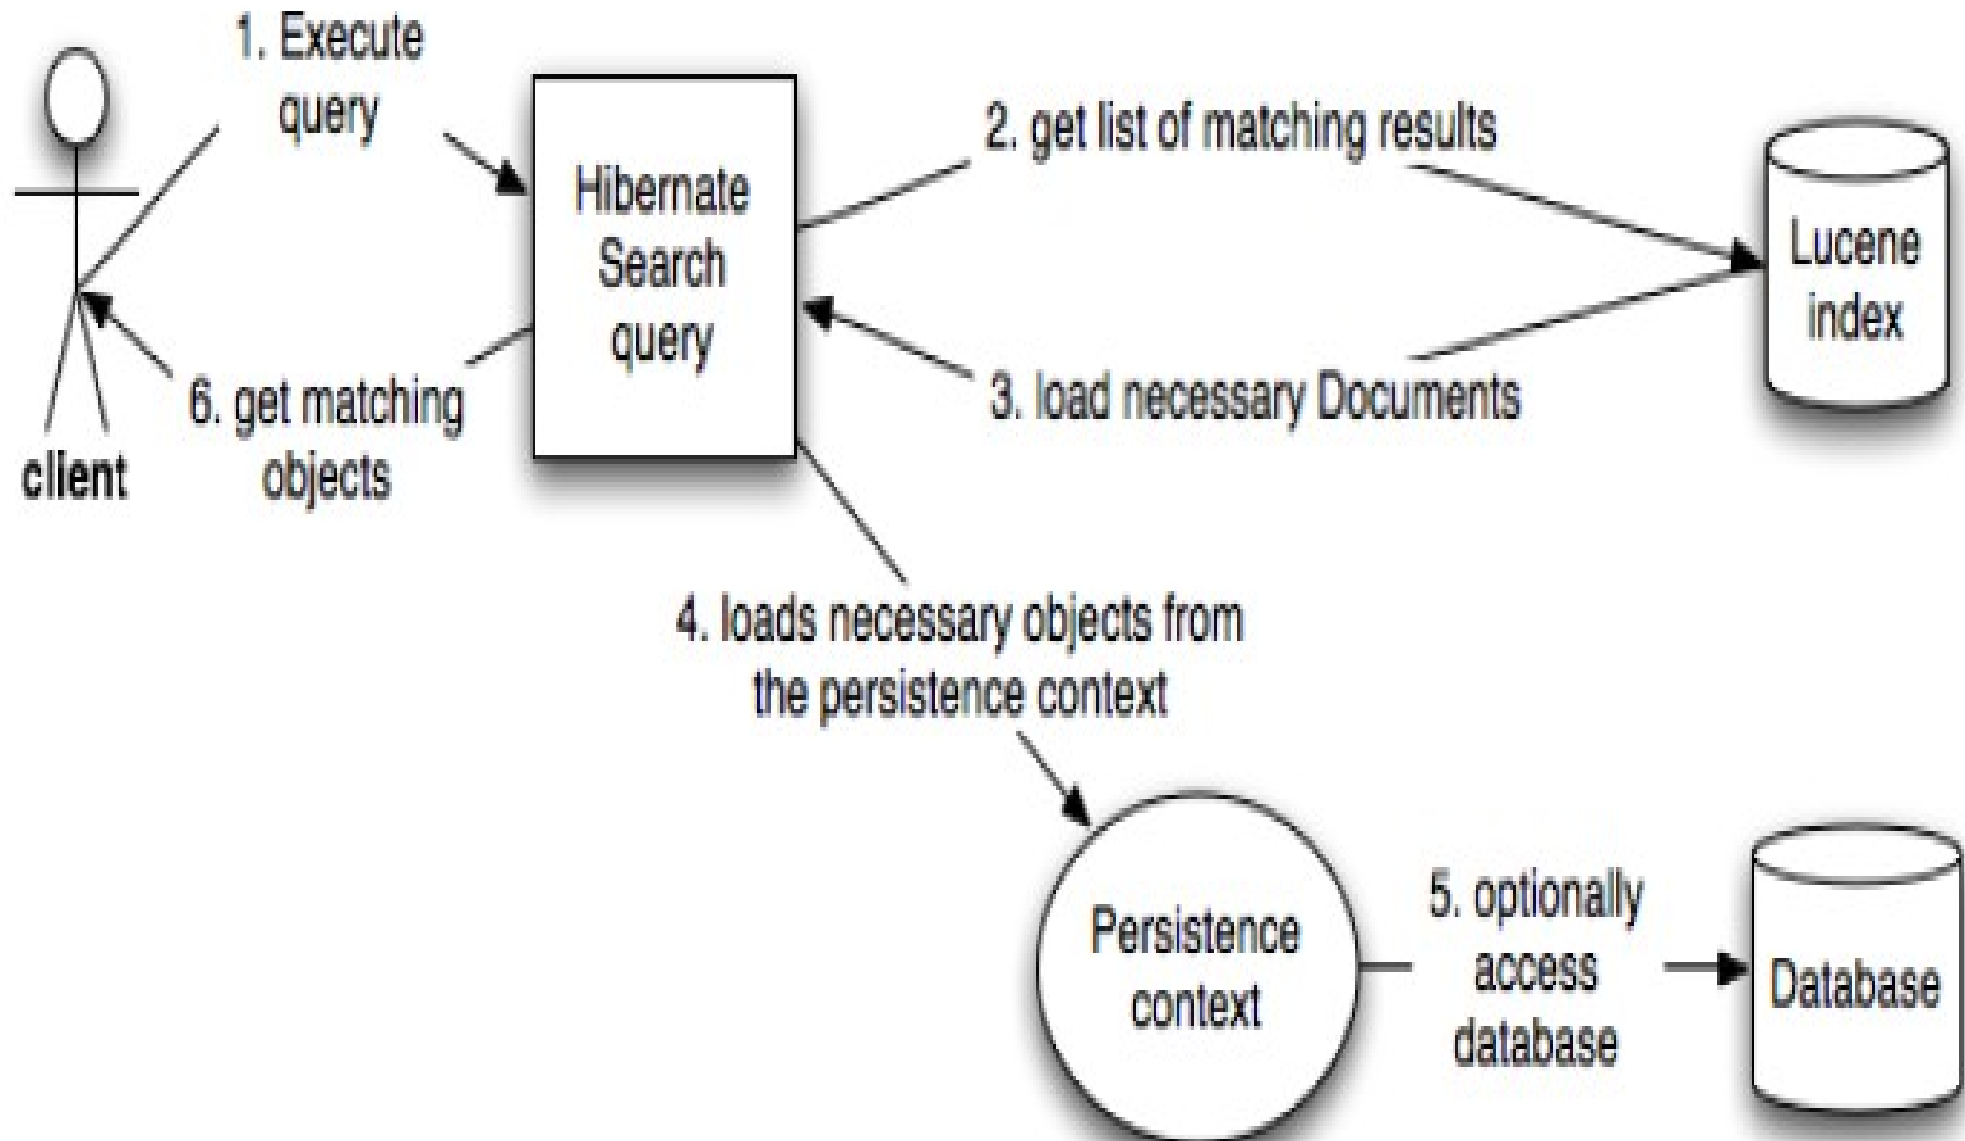
private FullTextEntityManager getFTEntityManager() {
    return Search.getFullTextEntityManager(em);
}

...
```

```
public List<Item> findByTitle(String words) {
    org.apache.lucene.search.Query luceneQuery =
        buildLuceneQuery(words, Item.class);
    javax.persistence.Query query =
        getFTEntityManager().createFullTextQuery(luceneQuery, Item.class);

    return query.getResultList();
}
```

Execute FullTextQuery



Pagination & Result Size

```
public Page<Item> search(String words,int pageNumber,int window) {
    org.apache.lucene.search.Query luceneQuery =
        buildLuceneQuery(words,Item.class);

    FullTextQuery query =
        getFTEntityManager().createFullTextQuery(luceneQuery,Item.class);

    List<Item> results = query
        .setFirstResult( (pageNumber - 1) * window )
        .setMaxResults(window)
        .getResultList();

    int resultSize = query.getResultSize();

    Page<Item> page = new Page<Item>(resultSize, results);
    return page;
}
```


Override Fetch Strategy

```
public List<Item> findByTitle(String words) {
    org.apache.lucene.search.Query luceneQuery =
        buildLuceneQuery(words, Item.class);

    FullTextQuery query =
        getFTSession().createFullTextQuery(luceneQuery, Item.class);

    Criteria fetchingStrategy =
        getFTSession().createCriteria(Item.class)
            .setFetchMode("actors", FetchMode.JOIN);

    query.setCriteriaQuery(fetchingStrategy);

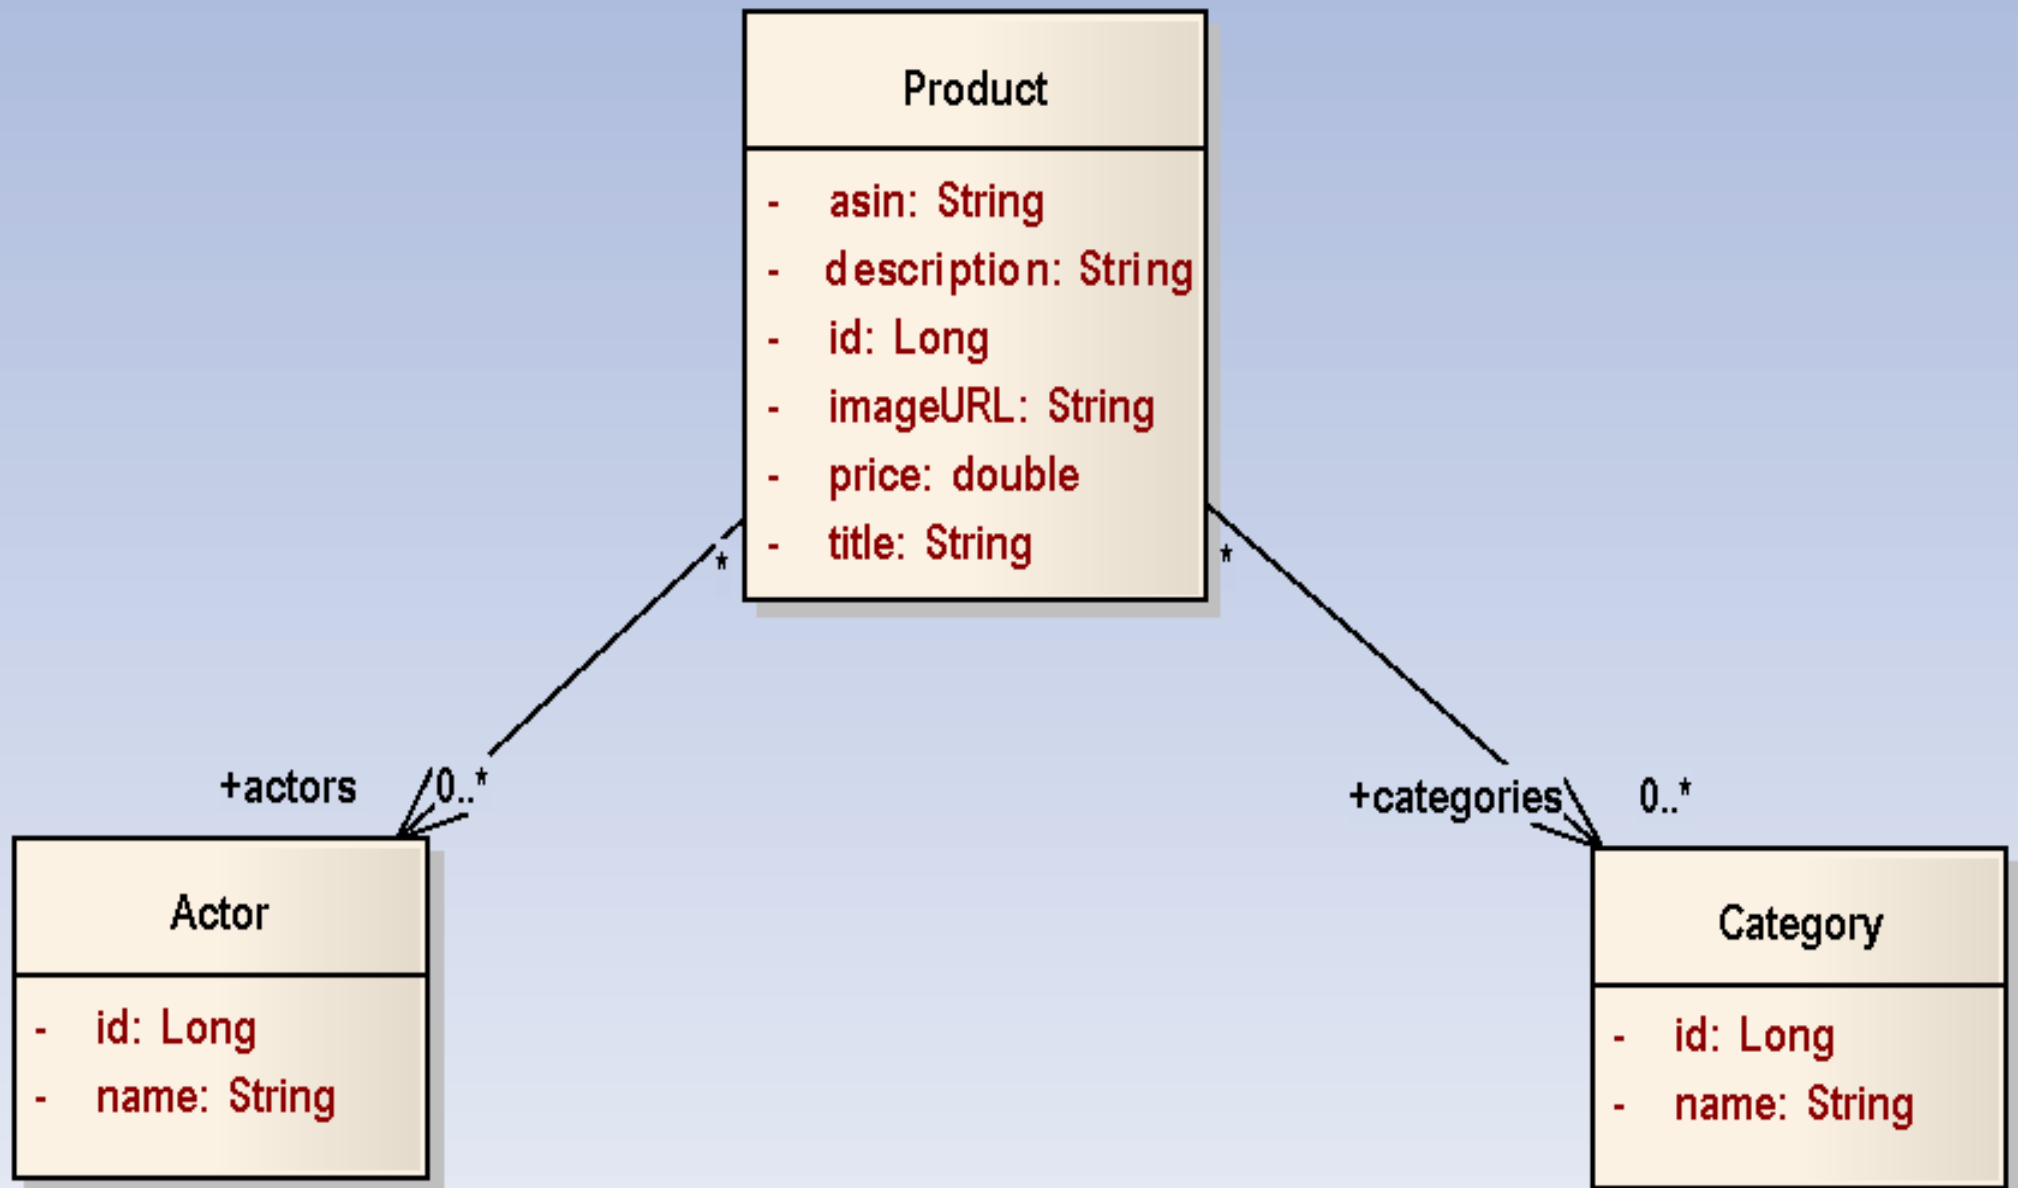
    return query.list();
}
```

- ▶ For JPA use getDelegate()
- ▶ Don't use Criteria restrictions
 - » Will hurt pagination and will provide wrong resultSize

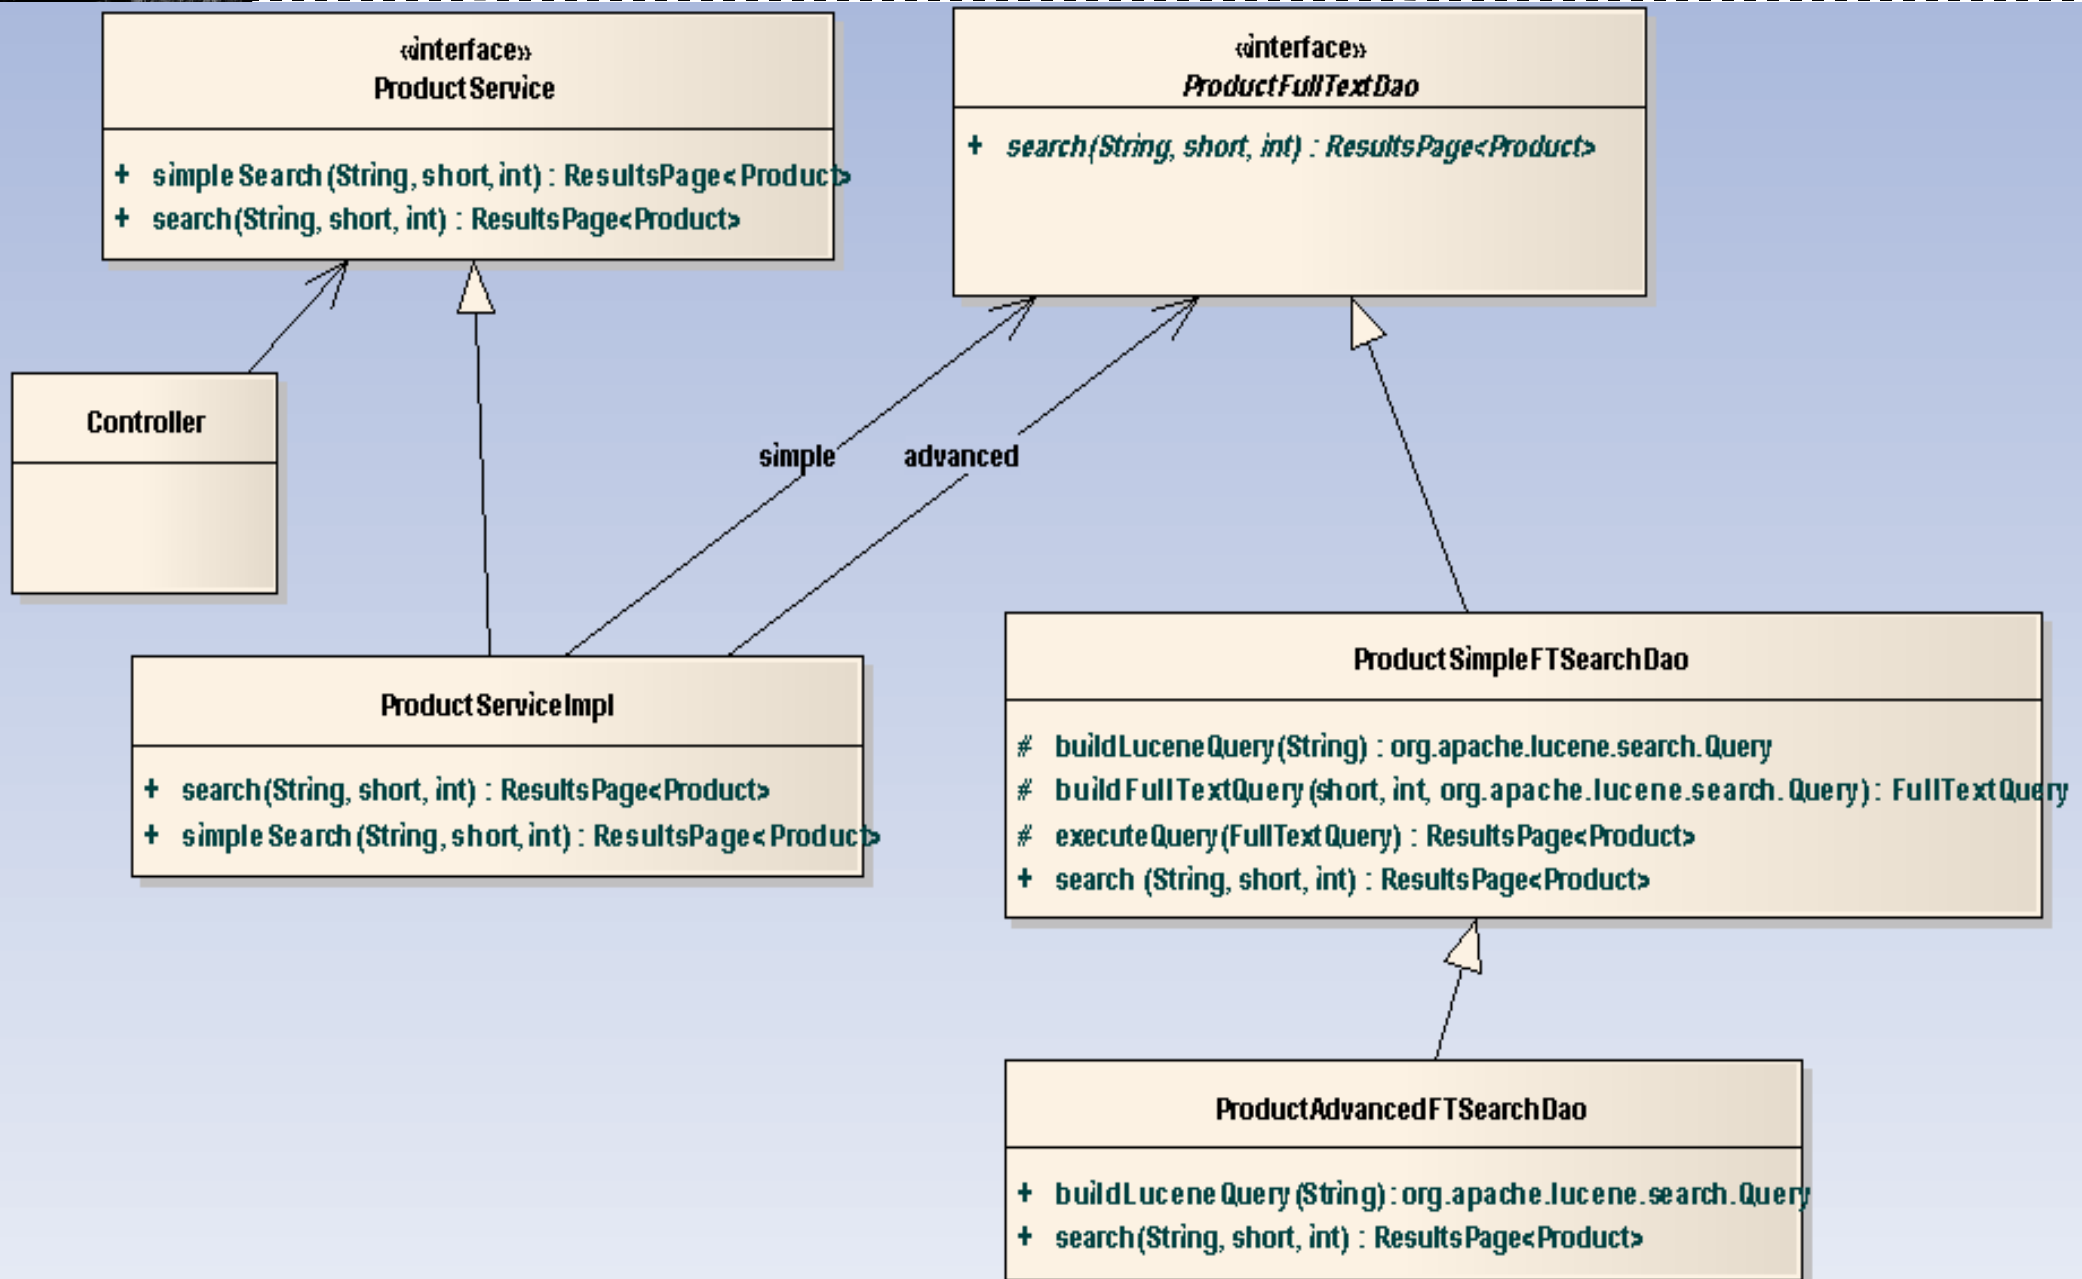
Demo



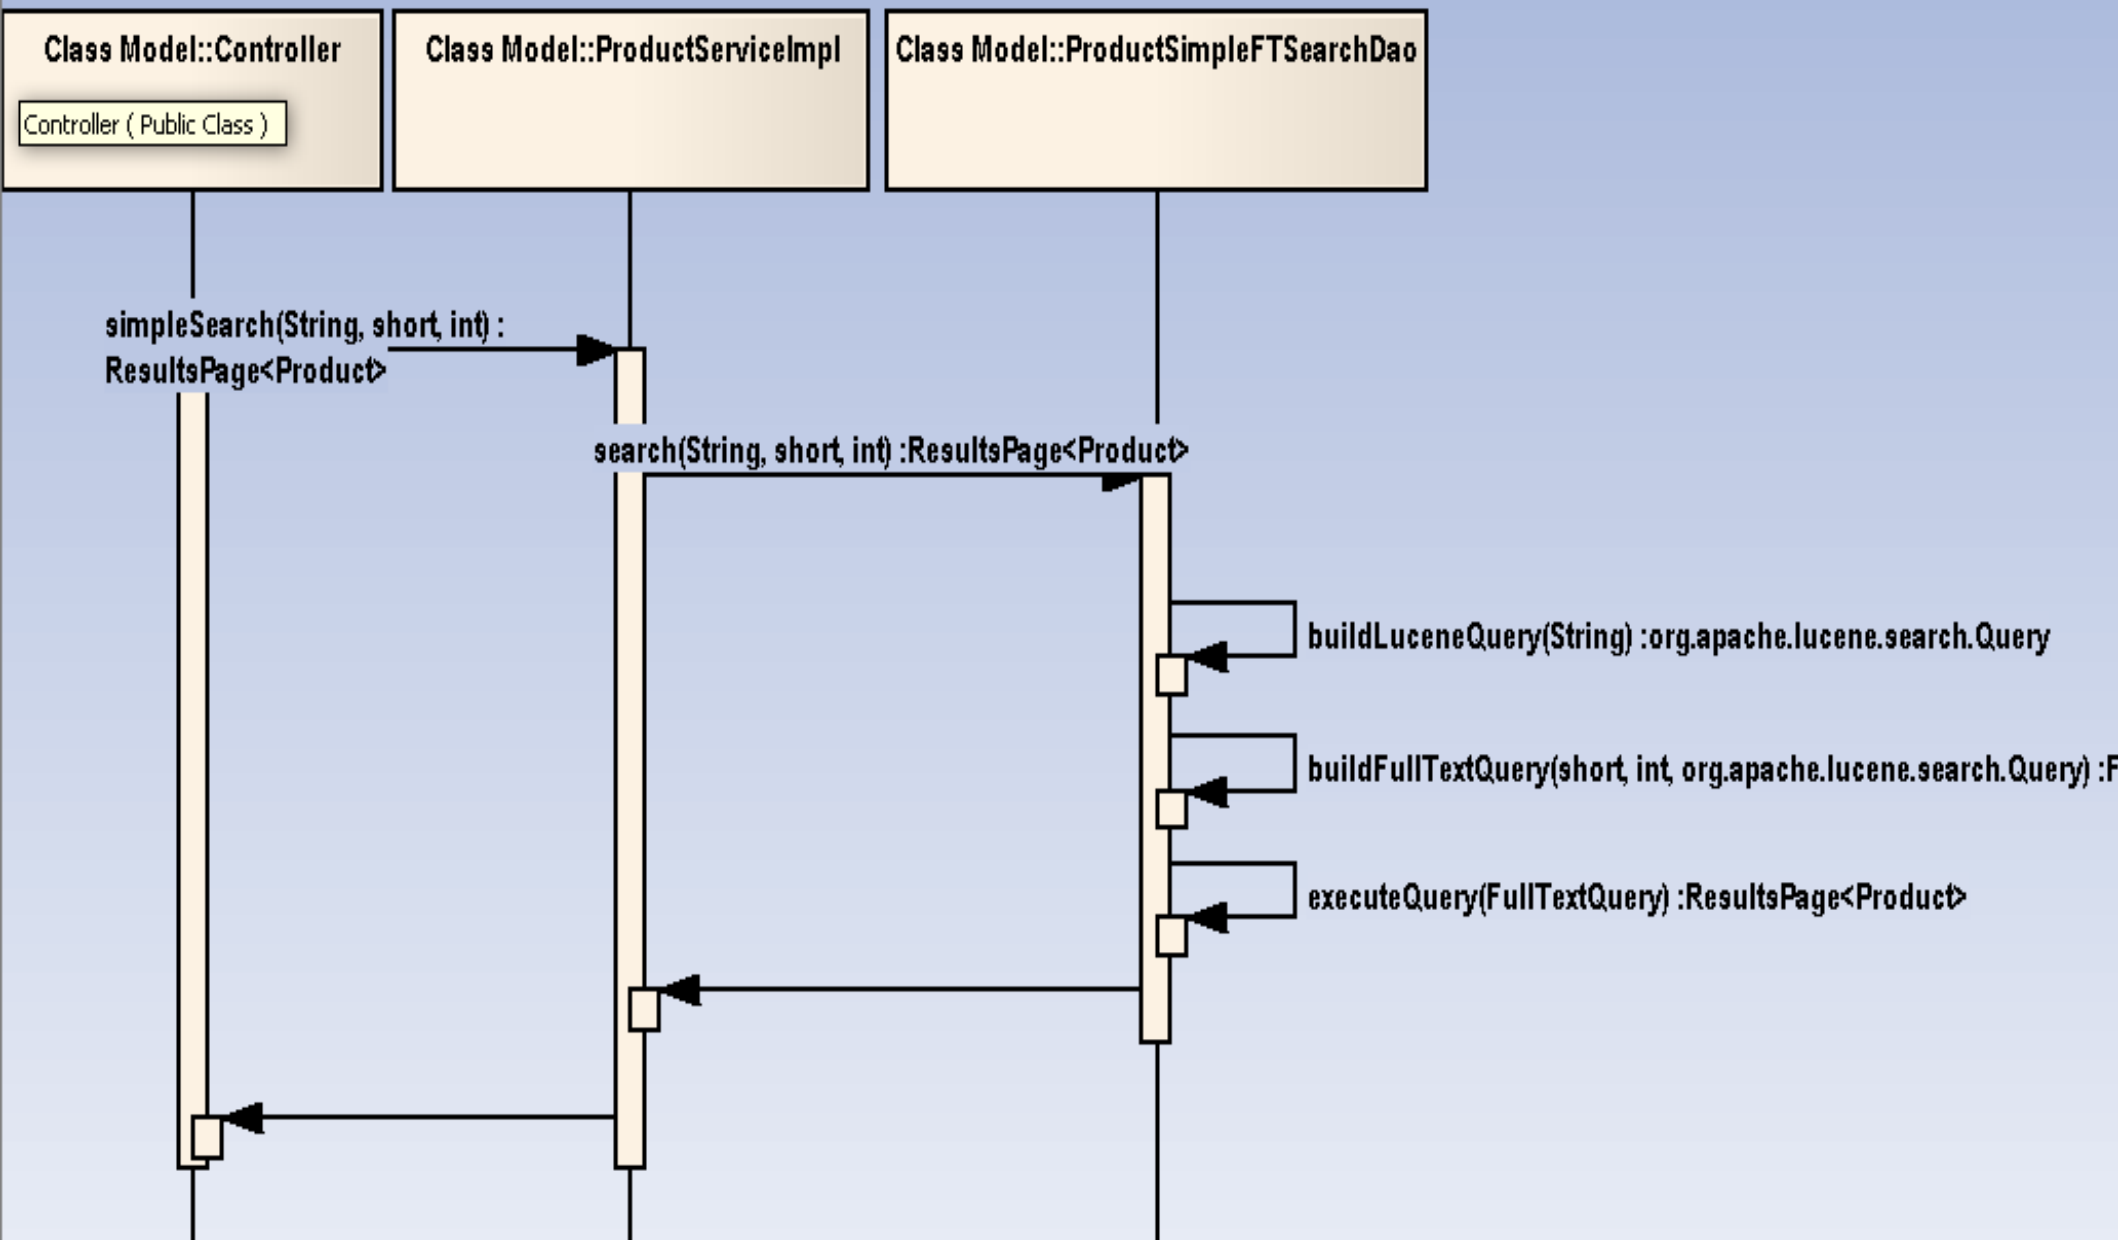
Product Domain Model

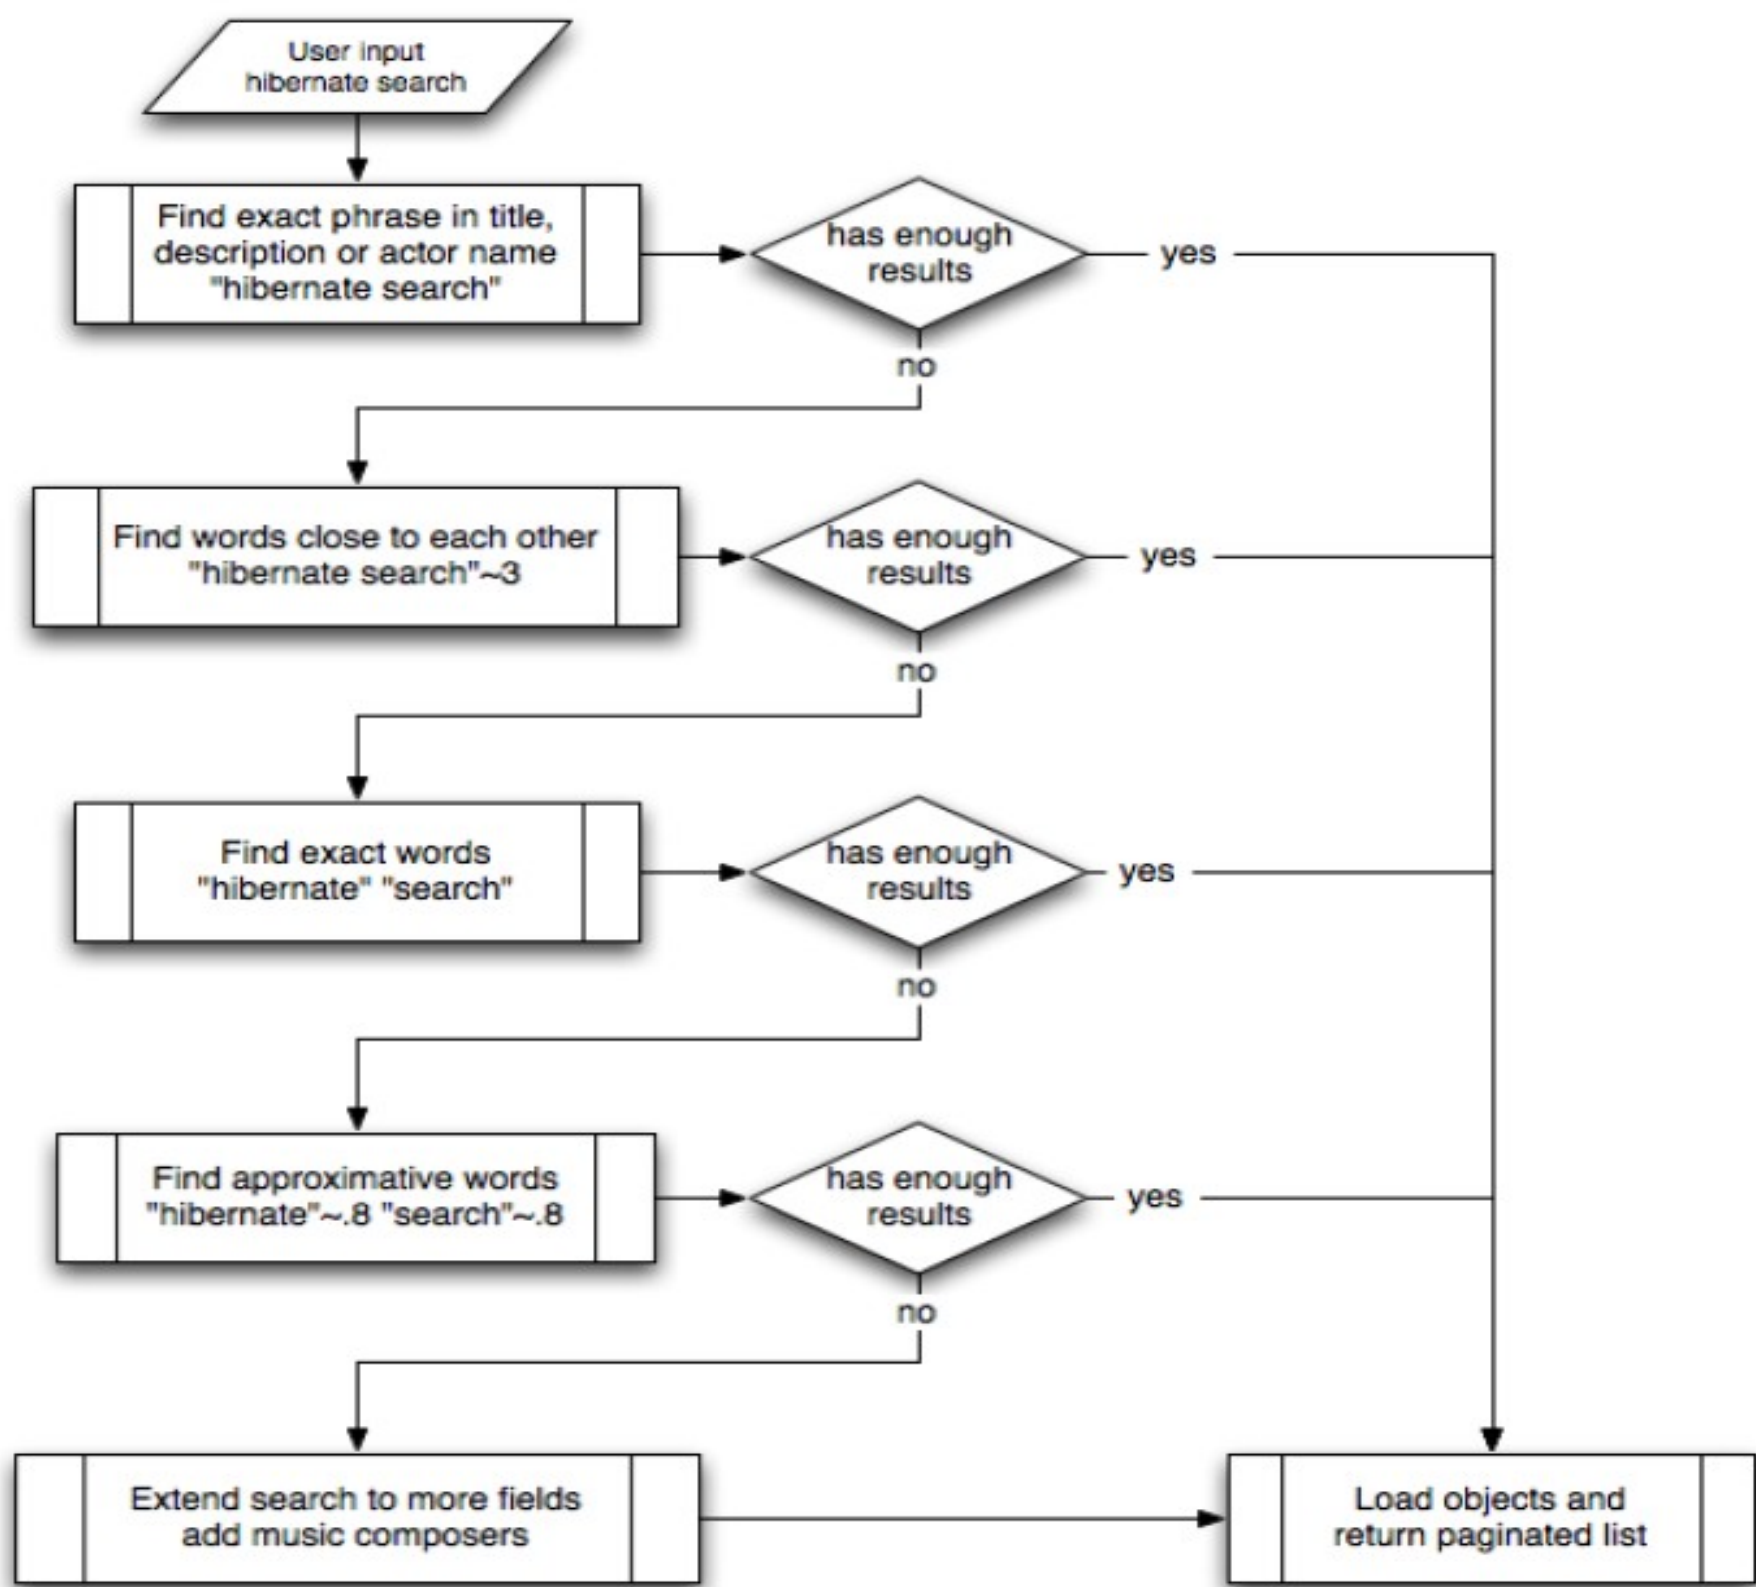


Service & DAO Layers



Simple Search Sequence Diagram





Projection

```
public class ItemView { // A view Object NOT necessary an entity...
    private String ean;
    private String title;
    public String getEan() {
        return ean;
    }
    public String getTitle() {
        return title;
    }
}
```

```
public List<ItemView> search(String words) {
    org.apache.lucene.search.Query luceneQuery =
        buildLuceneQuery(words, Item.class);
    FullTextQuery query =
        getFTSession().createFullTextQuery(luceneQuery, Item.class);
    query.setProjection("ean", "title");

    List<ItemView> results = query.setResultTransformer(
        new AliasToBeanResultTransformer(ItemView.class)).list();
    return results;
}
```

No hit on
the DB

Store Properties In The Index For Projection

```
@Entity @Indexed
public class Item {
    @Id @GeneratedValue
    private Integer id;

    @Field(store=Store.YES)
    private String title;

    @Field
    private String description;

    @Field(index=Index.UN_TOKENIZED, store=Store.YES)
    private String ean;
    ...
}
```


Sorting By Field

```
public List<Item> findByTitle(String words) {
    org.apache.lucene.search.Query luceneQuery =
        buildLuceneQuery(words, Item.class);

    FullTextQuery query =
        getFTSession().createFullTextQuery(luceneQuery, Item.class);

    Sort sort = new Sort(new SortField("title_sort", SortField.STRING));

    query.setSort(sort);

    return query.list();
}
```


```
@Entity @Indexed public class Item {
    ...
    @Fields({
        @Field(index=Index.TOKENIZED)
        @Field(name="title_sort", index=Index.UN_TOKENIZED)
    })
    private String title;
}
```

Dynamic Data Filtering

- ▶ Restrict results of a query **after the Lucene query has been executed**
 - » Rules that are not directly related to the query.
 - » Cross-cutting restrictions
 - category, availability , security.
- ▶ The ordering defined by the original query is respected.



Filters



Filter Example

```
//service implementation
public List<Item> searchItems(String search, boolean isChild) {
    org.apache.lucene.search.Query luceneQuery =
                                                buildLuceneQuery(search);

    FullTextQuery query =
        getFTSession().createFullTextQuery(luceneQuery, Item.class);

    if (isChild)
        query.enableFullTextFilter("chldFilter");

    List<Item> results = query.list();
    return results;
}
```

Filter Example Cont.

```
@Entity @Indexed
@FullTextFilterDef(name="childFilter",
                   impl=ChildFilterFactory.class)
@ClassBridge(name="childrenOnly",
              impl=ChildrenFlagBridge.class, index=Index.UN_TOKENIZED)
public class Item {
    ...
}
```

```
public class ChildFilterFactory {
    @Factory
    public Filter getChildrenFilter() {
        Query query = new TermQuery( new Term("childrenOnly", "yes") );
        return new QueryWrapperFilter( query );
    }
}
```

Optimizing Search

- ▶ Limit targeted classes (one class is the best)
 - » `ftSession.createFullTextQuery(luceneQuery, Item.class);`
- ▶ Use pagination
- ▶ Avoid the $n+1$ by using **`setCriteria()`**
- ▶ Use projection carefully






Scale Hibernate Search

Synchronous Clustering

- ▶ Who can use it?
 - » Applications with medium-size indexes
 - Network traffic will be needed to retrieve the index.
 - » Applications with low to moderate write intensive .



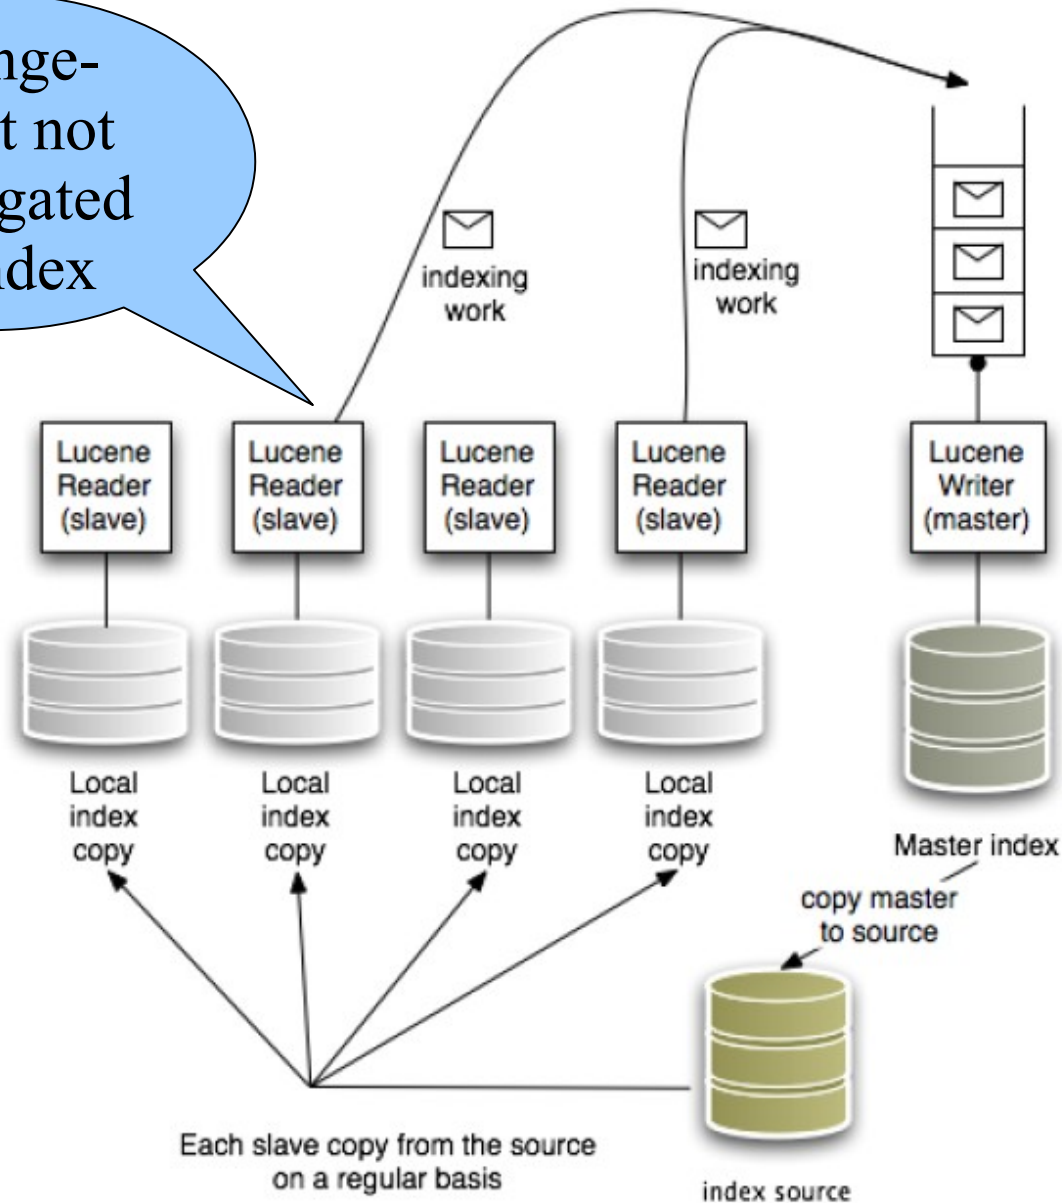
Synchronous Clustering Problems

- ▶ Some NFS cache the directory contents
 - » No immediate visibility for the directory content
 - Lucene relies (partially) on an accurate listing of files.
 - » “delete on last close” semantic NOT always implemented .
- ▶ Database Directory issues
 - » Segments are represented as blobs
 - » A pessimistic lock hurts concurrency on massive updates.
- ▶ In-memory distributed Directory
 - » GigaSpace, JBoss Cache and Terracotta



Asynchronous Clustering

Change-
Event not
propagated
to Index



Slave Configuration

```
<persistence-unit name="dvdstore-catalog">
  <jta-data-source>java:/DefaultDS</jta-data-source>
  <properties>
    <!-- regular Hibernate Core configuration -->
    <property name="hibernate.dialect"
              value="org.hibernate.dialect.H2Dialect"/>

    <!-- JMS backend -->
    <property name="hibernate.search.worker.backend"
              value="jms"/>
    <property name="hibernate.search.worker.jms.connection_factory"
              value="/ConnectionFactory"/>
    <property name="hibernate.search.worker.jndi.url"
              value="jnp://master:1099"/>
    <property name="hibernate.search.worker.jms.queue"
              value="queue/hibernatesearch"/>
  </properties>
</persistence-unit>
```

...

Slave Configuration Cont.

...

```
<!-- DirectoryProvider configuration -->
<property name="hibernate.search.default.directory_provider"
  value="org.hibernate.search.store.FSSlaveDirectoryProvider"/>
<property name="hibernate.search.default.refresh"
  value="1800"/>
<property name="hibernate.search.default.indexBase"
  value="/Users/prod/lucenedirs"/>
<property name="hibernate.search.default.sourceBase"
  value="/mnt/share"/>
</properties>
</persistence-unit>
```

Master Configuration

```
<persistence-unit name="dvdstore-catalog">
  <jta-data-source>java:/DefaultDS</jta-data-source>
  <properties>
    <!-- regular Hibernate Core configuration -->
    <property name="hibernate.dialect"
              value="org.hibernate.dialect.H2Dialect"/>

    <!-- Hibernate Search configuration -->
    <!-- no backend configuration necessary -->
    <property name="hibernate.search.default.directory_provider"
              value="org.hibernate.search.store.FSMasterDirectoryProvider"/>

    <property name="hibernate.search.default.refresh"
              value="1800"/>

    <property name="hibernate.search.default.indexBase"
              value="/Users/prod/lucenedirs"/>

    <property name="hibernate.search.default.sourceBase"
              value="/mnt/share"/>

  </properties>
</persistence-unit>
```

Building The Master MDB

```
@MessageDriven(activationConfig = {
    ActivationConfigProperty(propertyName="destinationType",
                                propertyValue="javax.jms.Queue"),
    @ActivationConfigProperty(propertyName="destination",
                                propertyValue="queue/hibernatesearch")
} )
public class MDBSearchController extends
    AbstractJMSHibernateSearchController implements MessageListener{

    @PersistenceContext private EntityManager em;

    @Override
    protected void cleanSessionIfNeeded(Session session) {
        //clean the session if needed nothing to do container managed
    }

    @Override
    protected Session getSession() {
        return (Session) em.getDelegate();
    }
}
```


What Happens On Master Failure ?

▶ Slave

- » Continue to serve full-text queries
- » Continue push changes that need indexing.

▶ Master

- » Messages on the master are roll-backed to queue.
- » Optional – Prepare a standby for the master

▶ On corrupted Index...

- » Re-index manually from DB
- » Optional – Use Storage Area Network (SAN)



Summary



Full-Text Search Without The Hassle

- ▶ Solves The 3 mismatch problems
 - » Automatic **structural conversion** through Mapping
 - » Transparent **index synchronization**
 - » **Retrieved data** from index become “**persistent**” entities.
- ▶ **Easier** / Transparent optimized Lucene use
- ▶ **Scalability** capabilities out of the box





Q & A




Thank You

yanai@tikalk.com



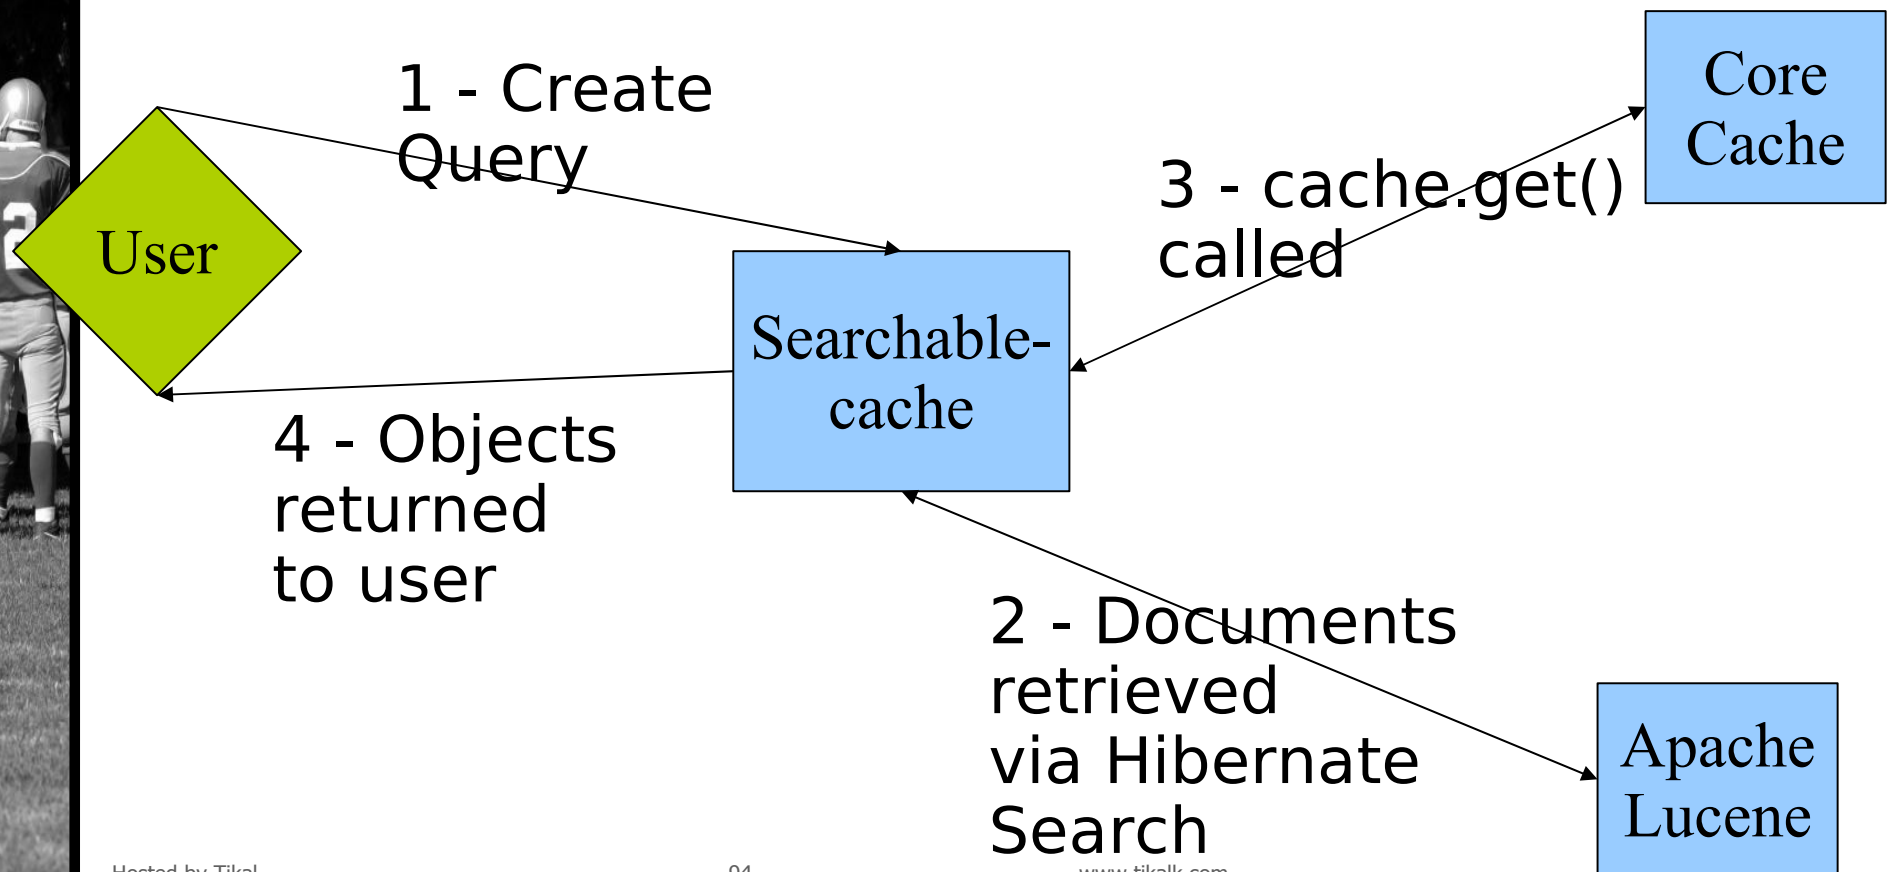
Appendixes

Use SAN For Lucene Directory



JBoss Cache Searchable

- ▶ Integration package between **JBoss Cache** and **Hibernate Search**.
- ▶ Provides **full text search** capabilities to the **cache**.



Annotated Pojo

```
@Indexed
@ProvidedId
public class Person { //Not necessary a Hibernate Entity

    @Field
    private String name;

    @Field
    private Date dateOfBirth;

    //Not Indexed
    private String massiveString;

    //Standard getters, setters etc follow.

}
```

FullText Search on Cache

```
Cache<String, Person> cache =  
    new DefaultCacheFactory<String, Person>().createCache();  
SearchableCache searchableCache = new SearchableCacheFactory();  
    createSearchableCache(cache, Person.class);  
...
```

```
public void putStuffIn(Person p){  
    searchableCache.put(Fqn.fromString("/a/b/c"), p.getName(), p);  
}
```

```
public List findStuff(String searchStr){  
    Query luceneQuery = buildLuceneQuery(searchStr)  
  
    CacheQuery cacheQuery =  
        searchableCache.createQuery(luceneQuery, Person.class);  
  
    return cacheQuery.list();  
}
```