
JBoss and Spring: Integration Features

By : Zvika Markfeld, Tikal



Agenda

- ▶ Introduction
- ▶ JBoss Spring Deployer
- ▶ JBossMX and Spring JMX Support
- ▶ Spring Clustering with JBoss Cache
- ▶ Conclusion
- ▶ Resources





Introduction

Introduction to Spring

- ▶ Over the several several years, the Spring Framework has emerged as the de facto standard for developing **simple, flexible, easy-to-configure** Enterprise applications
- ▶ At the heart of Spring lies the ***Inversion of Control*** (IoC) design pattern:
 - » Application as a set of POJOs
 - » IoC Container wires them together, setting dependencies

Introduction to Spring – Cont'

- ▶ The Spring container is configured via a set of bean definitions, typically expressed via XML context files:

```
<bean id="MyServiceBean" class="mypackage.MyServiceImpl">  
  <property name="otherService" ref="OtherServiceBean"/>  
</bean>
```

- ▶ Client Code:

```
MyServiceInterface service =  
    (MyServiceInterface)context.getBean("MyServiceBean");  
service.doSomething();
```

Introduction to Spring – Cont'

- ▶ In addition to IoC, Spring provides many other services, coding conveniences, and "hooks" into standard APIs that ease server-side development
 - » EJB(2.1, 3.0)
 - » JMS
 - » JMX
 - » MVC frameworks
 - » Transparent Remoting
 - » And more...

Spring and EJB3

	EJB3	Spring
Dependency Injection	Coarse-grained, simple, JNDI-based	Fine-grained, versatile instantiation/lookup
Development Cycle	code -> compile -> package -> deploy -> test	short, environment-agnostic cycle
Transaction Propagation	In and across-Containers	In-Container only
AOP	Simple interceptor support	Full AOP stack
Transaction Manager	Container-provided	Pluggable
Persistence	JPA	JPA, Hibernate, JDO, etc.
Clustering	Non-mandatory, usually provided	Not provided, roll-your-own

Spring and EJB3 Integration

- ▶ Spring DAO Support:
 - » JpaTemplate, JpaDaoSupport, JpaTransactionManager, etc.
 - » EntityManager Injection via PostProcessor
 - » Exception Translation
- ▶ Coarse-grained EJBs -> Fine grained Spring Services
 - » EJB Container provides:
 - Clustering
 - Transaction propagation across JVMs
 - Classloading domains
 - » Spring provides:
 - Everything else...
- ▶ The Pitchfork Project
 - » WLS and Interface21 initiative, open source
 - » Pluggable EJB3 container on top of Spring
 - » Implements JSR220, JSR250

Spring Remoting Example

- ▶ A simple approach to Java-to-Java remoting in Spring is **HTTP Remoting**
- ▶ For example, the following context piece exposes MyService for public consumption (after servlet registration in web.xml):

```
<bean name="/MyRemoteService"
      class="org.springframework...HttpInvokerServiceExporter">
  <property name="service" ref="MyServiceBean"/>
  <property name="serviceInterface"
    value="mypackage.MyServiceInterface"/>
</bean>
```

- ▶ Actual service is injected into this bean definition and thus made available for the remote calls

Spring Remoting Example – Cont'

- ▶ On the client, the context definition reads:

```
<bean id="myServiceBean"
      class="org.springframework...HttpInvokerProxyFactoryBean">
  <property
    name="serviceUrl"
    value="http://somehost:8080/my-webapp/MyRemoteService"/>
  <property
    name="serviceInterface"
    value="mypackage.MyServiceInterface" />
</bean>
```

- ▶ Client-side code doesn't change!

```
MyServiceInterface service =
    (MyServiceInterface)context.getBean("MyServiceBean");
service.doSomething();
```

Spring Remoting – Up and Away...

- ▶ Besides HTTP, Spring supports exporting POJOs over several other remoting protocols, out of the box:
 - » Web services
 - » Hessian
 - » Burlap
 - » RMI
 - » JMX(discussed later)
 - » And others...



Spring and JBoss – a Match Made In Heaven?

- ▶ Apart for the ‘political’ disputes, it seems that there’s a good match:
 - » Advanced, innovative technology providers
 - » Emphasize extensibility(each in its own domain)
 - » Focus on the Enterprise Applications domain
 - » Similar open source licensing
- ▶ Spring starts where JBoss stops:
 - » Abstraction layers for common services
 - » Standard, developer-friendly programming model
- ▶ In the next slides, we will examine some areas of integration...



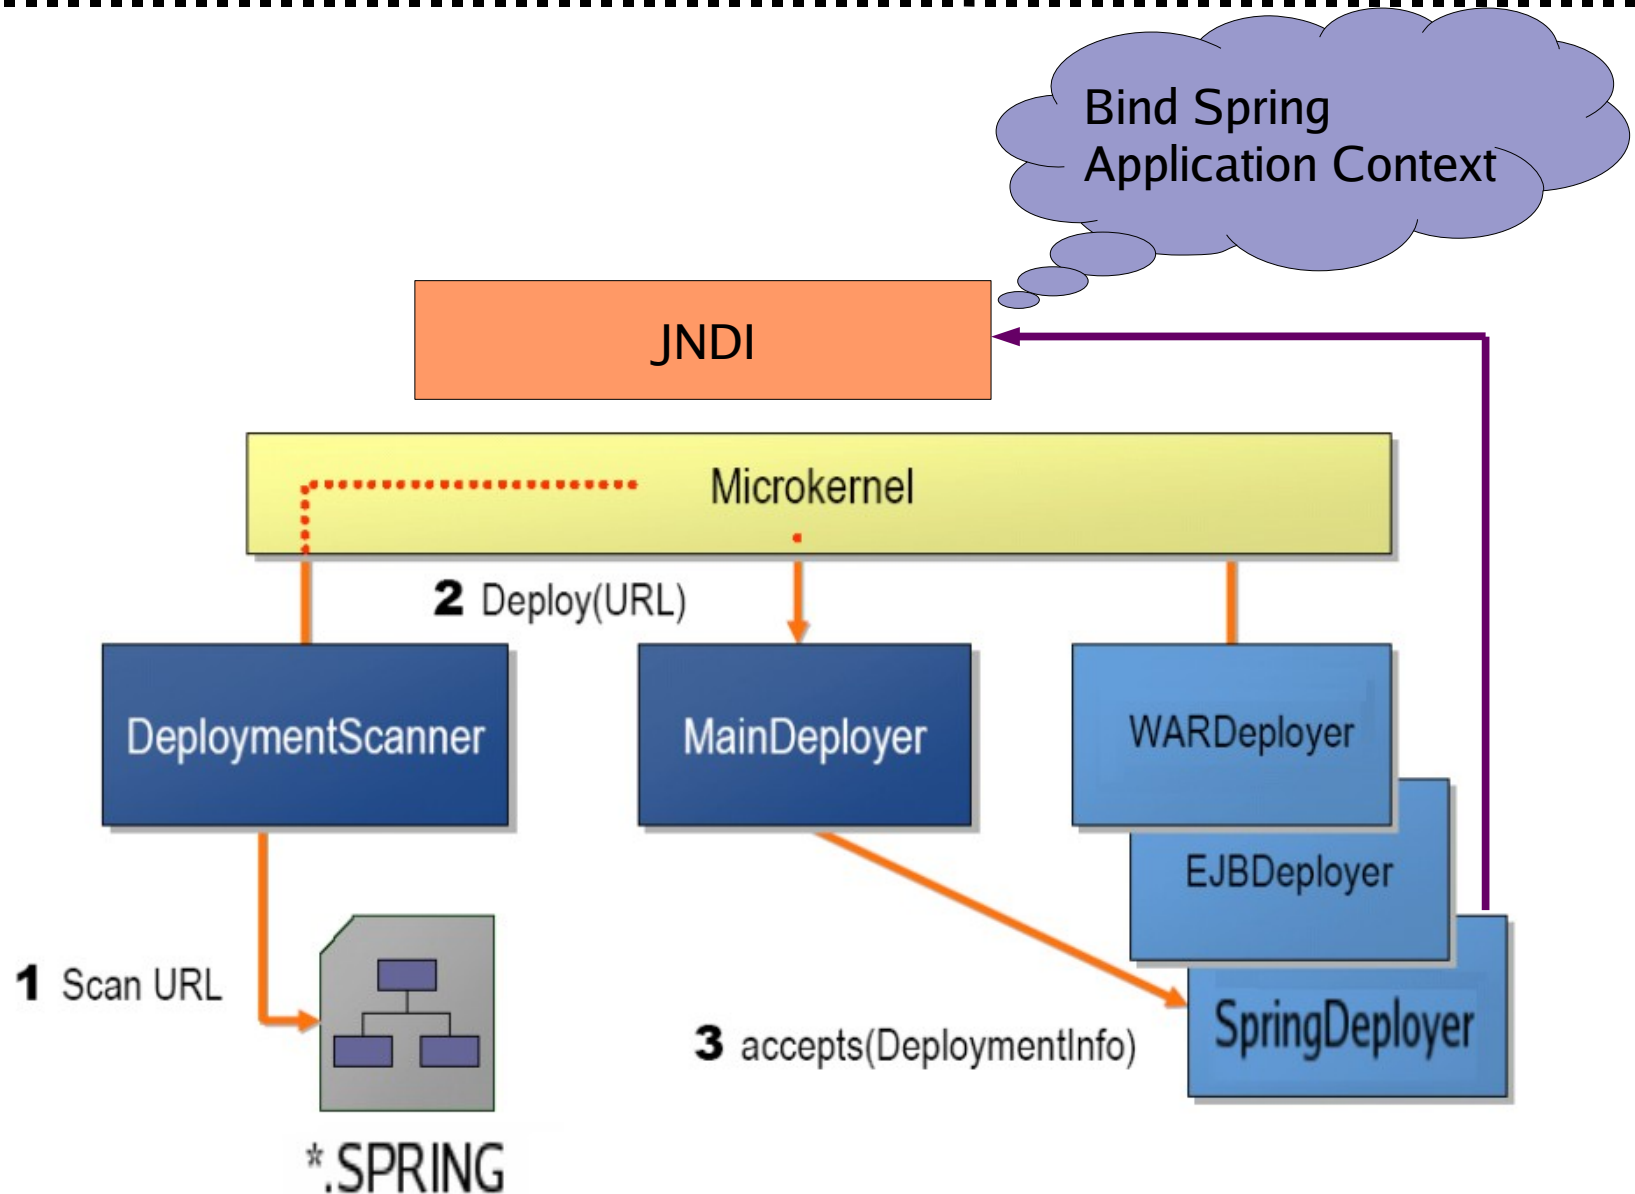


JBoss Spring Deployer

JBoss Spring Deployer

- ▶ Allows deploying *.spring archives into JBoss
 - » Create JAR archive with a **META-INF/jboss-spring.xml**
- ▶ EJB 3.0 integration also supported
 - » Inject Spring-declared beans into an EJB using a **@Spring** annotation
 - » Based on JBoss 4's AOP capabilities
 - » Distributed as a JBoss 4 deployer archive
- ▶ Written by Ales Justin of jboss.org...

JBoss Spring Deployments Architecture



Spring Deployment Structure

my-app.spring/

org/

acme/

~MyBean.class

~MyBean2.class

META-INF/

jboss-spring.xml

Bean Classes

Standard Spring
Configuration

Spring Annotations Example

```
@Stateless
@Interceptors(SpringLifecycleInterceptor.class)
public class RandomizerBean implements Randomizer {

    private WordsCreator wordsCreator;

    @Spring(jndiName = "my-app", bean="stateIntCreator")
    private IntCreator intCreator;

    public WordsCreator getWordsCreator() {
        return wordsCreator;
    }

    @Spring(jndiName = "my-app", bean="staticWordsCreator")
    public void setWordsCreator(WordsCreator wordsCreator) {
        this.wordsCreator = wordsCreator;
    }
}
```

Spring Archives Deployment

- ▶ The Deployer will register the XML defined bean factory into JNDI, in a non-serialized form
- ▶ Default JNDI name = short deployment filename (**my-app** in the example)
- ▶ Alternatively, drop the Spring XML file of the form **<name>-spring.xml** into the JBoss deploy directory
 - » Put your jar libraries under `server/<conf-name>/lib`
- ▶ You can also embed these deployments inside an EAR, EJB-JAR, SAR, etc. as nested archives

Spring Deployment Metadata

- ▶ JNDI Naming and context inheritance can be marked in jboss-spring.xml or *-spring.xml, as follows:

```
<beans>
  <description>
    BeanFactory=( <name> )
    ParentBeanFactory=( <parent name> )
  </description>
  ...
</beans>
```



JBossMX and Spring JMX Support

JBossMX and Spring JMX Support

- ▶ Spring JMX support provides 4 core features for transparent JMX integration, applicable for JBossMX:
 - » Automatic Registration of any POJO as a JMX MBean
 - » Flexible mechanism for controlling the management interface of your beans
 - » Simple proxying of local and remote MBean resources
 - » Declarative exposure of MBeans over remote, JSR-160 connectors



Exporting Spring Beans Over JBossMX

```
<beans>
  <bean id="testBean" class="com.mycompany.MyTestBean">
    <property name="name" value="TEST"/>
    <property name="age" value="100"/>
  </bean>
  <bean id="exporter"
    class="org.springframework.jmx.export.MBeanExporter">
    <property name="beans">
      <map>
        <entry key="bean:name=testBean1" value-ref="testBean"/>
      </map>
    </property>
  </bean>
  ...
</beans>
```

Exporting Spring Beans Over JBossMX – Cont'

- ▶ *testBean* is exposed as a JMX MBean under the ObjectName **bean:name=testBean1**
 - » Default ObjectName can be overridden by supplying a custom NamingStrategy object to the Exporter
- ▶ Bean's public properties and methods are exposed as MBean attributes and operations
- ▶ Lazy bean instantiation will be respected by the MBeanExporter, using a proxy



Automatic MBean Registration

- ▶ MBeans can be automatically detected by the MBeanExporter by setting the *autodetect* property to true:

```
<bean id="exporter"  
      class="org.springframework.jmx.export.MBeanExporter">  
  <property name="autodetect" value="true"/>  
</bean>  
  
<bean name="spring:mbean=true"  
      class="com.mycompany.MyTestBean"/>
```

Bean should
implement an
interface
ending with
MBean

Fine-Tuning MBean Management Interface

- ▶ Behind the scenes, the MBeanExporter uses an *MBeanInfoAssembler* implementation for defining the exposed management interface of each bean
- ▶ The default implementation defines an interface that exposes all public properties and methods
- ▶ Spring provides two additional strategies for controlling the management interface:
 - » Using source level metadata
 - Commons Attributes
 - Java5 Annotations
 - » Using an arbitrary Java Interface



Using Java5 Annotation

```
@ManagedResource(objectName="bean:name=testBean4",
    log=true, logFile="jmx.log", currencyTimeLimit=15,
    persistPolicy="OnUpdate", persistPeriod=200,
    persistLocation="foo", persistName="bar")

public class AnnotationTestBean implements IJmxTestBean {
    private String name;

    @ManagedAttribute(currencyTimeLimit=20,
        defaultValue="bar", persistPolicy="OnUpdate")
    public void setName(String name) {
        this.name = name;
    }

    @ManagedOperation(description="Add Two Numbers Together")
    public int add(int x, int y) {
        return x + y;
    }
    ...
}
```


Annotation Configuration

```
<beans>
  <bean id="exporter"
    class="org.springframework.jmx.export.MBeanExporter">
    <property name="beans">
      ...
    </property>
    <property name="assembler" ref="assembler"/>
  </bean>
  <bean id="assembler"
    class="org.springframework.jmx...MetadataMBeanInfoAssembler">
    <property name="attributeSource" ref="attributeSource">
  </bean>

  <bean id="attributeSource"
    class="org.springframework.jmx...AnnotationJmxAttributeSource"/>
</beans>
```

Accessing MBeans via Proxies

- ▶ Spring JMX allows you to create proxies that re-route calls to MBeans registered in a local or remote MBeanServer
- ▶ These proxies provide you with a standard Java interface for interacting with MBeans

```
<bean id="proxy"
  class="org.springframework.jmx.access.MBeanProxyFactoryBean">

  <property
    name="objectName" value="bean:name=testBean" />

  <property
    name="proxyInterface" value="com.mycompany.IMyTestBean" />

</bean>
```



Spring Clustering with JBoss Cache

Spring Clustering with JBoss Cache

- ▶ **The domain:**
Advertising pojos over http invokers
- ▶ **The problem:**
URLs service identifiers are too brittle
 - » Server restarts, topology modification resistance
- ▶ **Possible solution:**
Employ a naming service to provide dynamic, real-time resolution of service names
 - » Keep-alive based approach for maintaining online destinations: slow, non-transactional
- ▶ **...Yet better solution:**
Distributed Cache
 - » Simple Map-like interface stores service endpoints

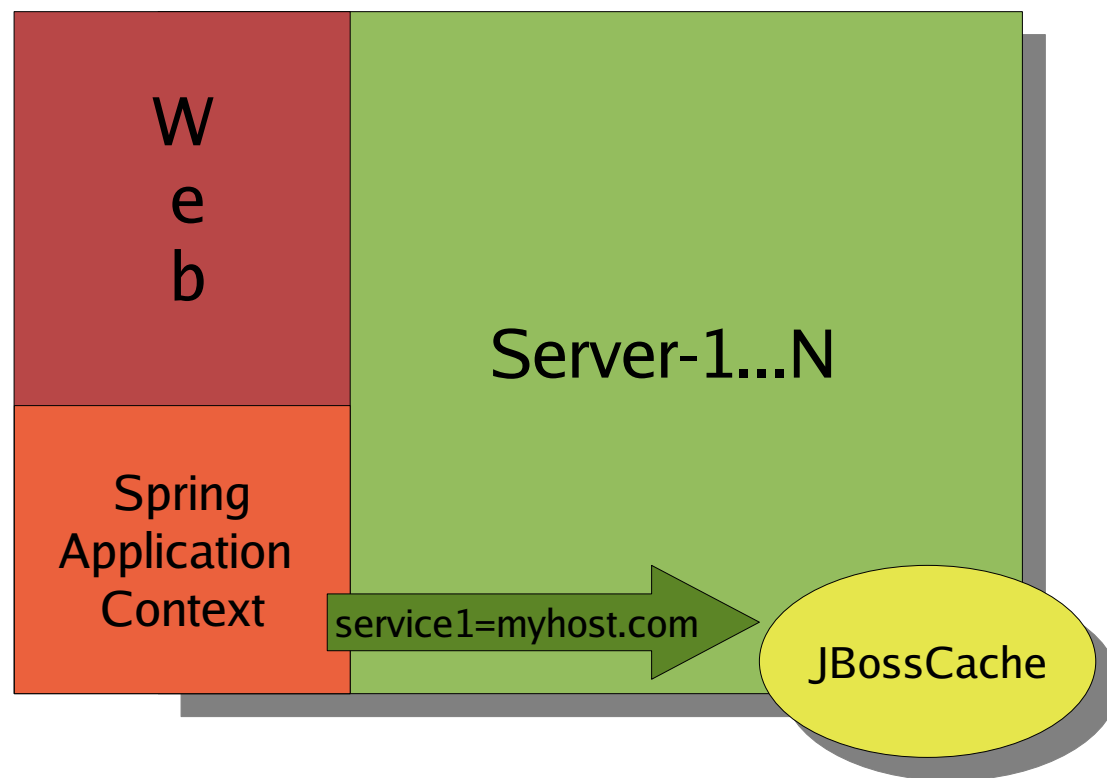
Distributed Cache as Dynamic Service Registry

- ▶ Associate a service name with one or more URLs pointing to its implementations
- ▶ Store the [name=(list of URLs)] associations in a distributed cache
- ▶ Update them as the network situation changes
- ▶ Service clients uses cache to inquire and access service implementations
- ▶ Bonus: client-side load balancing and failover
 - » Scalability better than in stateful session replication



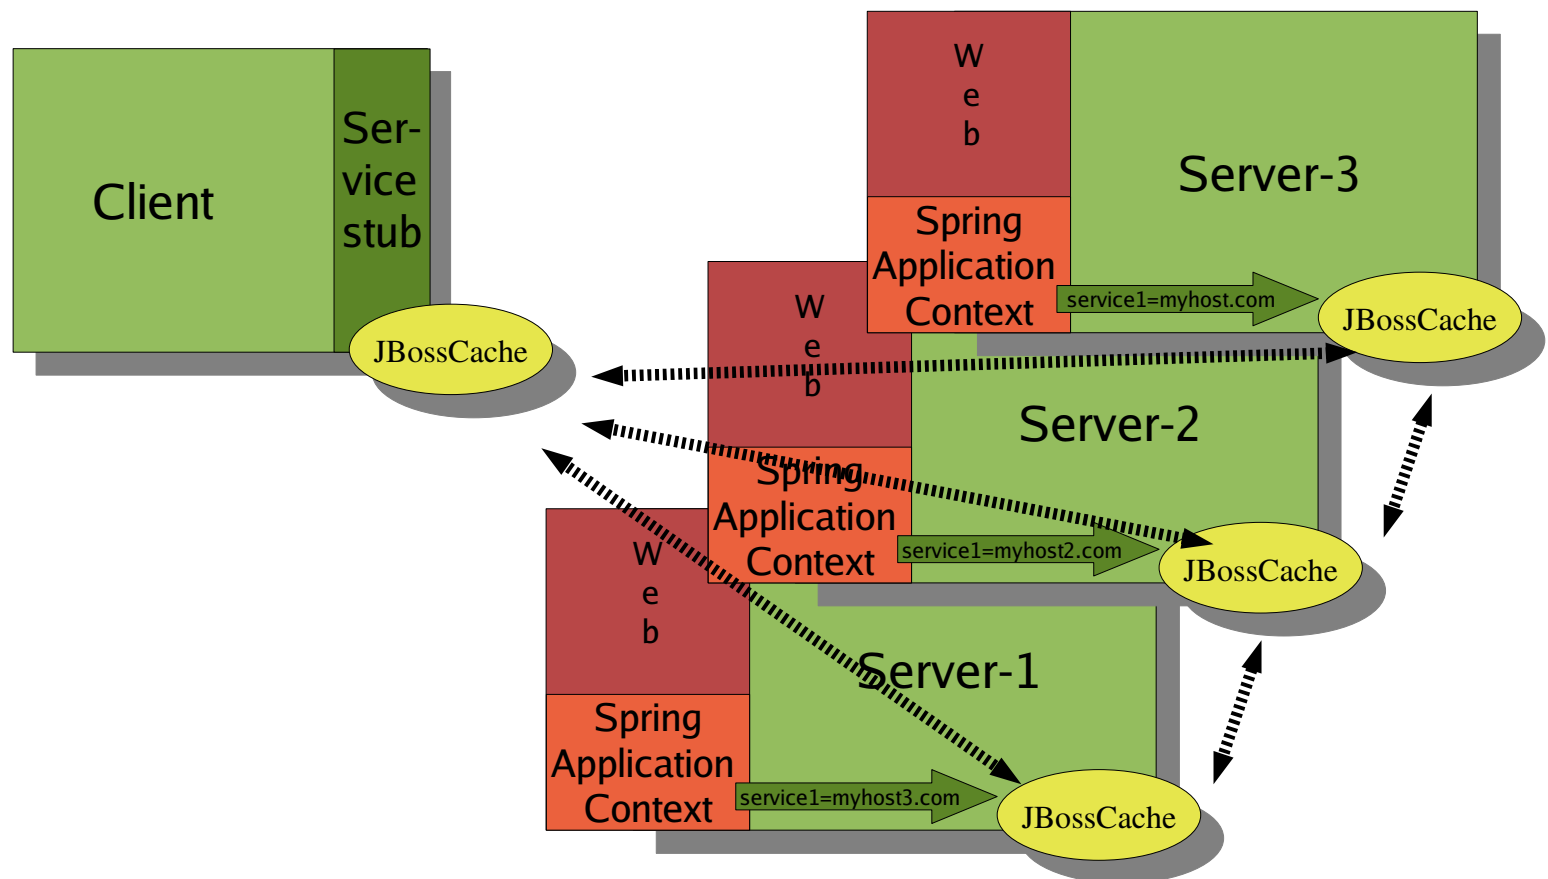
Suggested Topology

- ▶ Each server register the services it exports in the Cache, using the service's logical name as a key



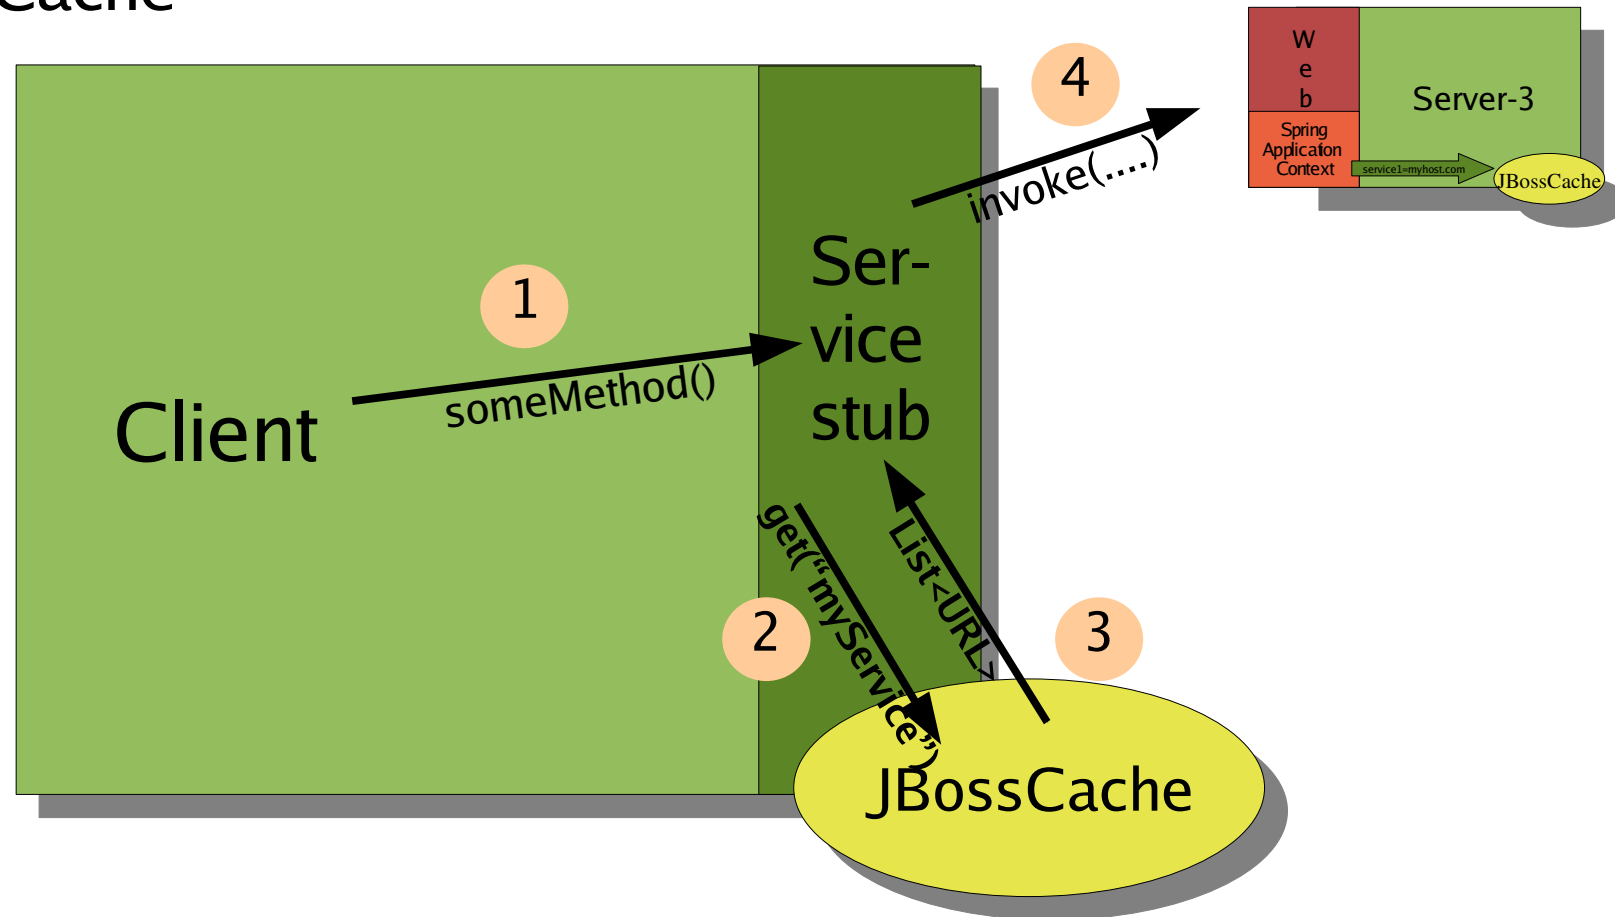
Suggested Topology – Cont'

- ▶ All participating nodes share the same cache, and so does the service client, using a cluster-aware stub



Suggested Topology – Cont'

- ▶ When the client needs to access the service, the smart stub finds an available endpoint by querying the local Cache



Setting Up Service Distribution

► Step 1: Exposing our TestService for remoting

```
<bean name="myService" class="com.mycompany.TestService"/>

<bean name="/TestService"
      class="org.springframework...HttpInvokerServiceExporter">
  <property name="service" ref="testService"/>
  <property name="serviceInterface"
    value="app.service.TestServiceInterface"/>
</bean>
```

Cache Proxy

- ▶ **Step 2: Creating a Server-Side JMX Proxy for the JBoss Cache MBean**

```
<bean id="customTreeCacheMBean"  
      class="org.springframework.jmx.access.MBeanProxyFactoryBean">  
  <property name="objectName">  
    <value>jboss.cache:service=CustomTreeCache</value>  
  </property>  
  <property name="proxyInterface">  
    <value>org.jboss.cache.TreeCacheMBean</value>  
  </property>  
</bean>
```

Cache Service

► Step 3: Writing an wrapper for the CacheMBean

```
public class JBossCacheServiceImpl
    implements CacheServiceInterface, InitializingBean {

    private TreeCacheMBean cacheMBean;

    public void put(String path, Object key, Object val)
        throws Exception {
        cacheMBean.put(path, key, value);
    }

    public Object get(String path, Object key) throws Exception {
        return cacheMBean.get(path, key);
    }

    public void setCacheMBean(TreeCacheMBean cacheMBean) {
        this.cacheMBean = cacheMBean;
    }

    public void afterPropertiesSet() throws Exception {
        cacheMBean.addTreeCacheListener(new MyCacheListener());
    }
}
```

Service Publisher

- ▶ **Step 4:** On each server node, a Spring Listener publishes and removes local services to/from Cache

```
private void contextRefreshed() throws Exception {
    logger.info("context refreshed");
    String[] names =
        context.getBeanNamesForType(HttpInvokerServiceExporter.class);
    logger.info("exporting services:" + names.length);
    for (int i = 0; i < names.length; i++) {
        String serviceUrl = makeUrl(names[i]);
        Set services = (Set)cache.get("/auto/svc/" + names[i], "KEY");
        if (services == null) {
            services = new HashSet();
        }
        services.add(serviceUrl);
        cache.put("/auto/svc/" + names[i], "KEY", services);
        logger.info("added:" + serviceUrl);
    }
}
```

Service Publisher Mechanics

- ▶ The publisher iterates over the list of services exported on http and adds their URLs to the cache
- ▶ The cache region, or path, contains the name of the service, whose URL list is stored as a `java.util.Set` under a fixed key in this region
- ▶ With JBoss Cache, it is important to make the service names part of the path to provide proper transactional scoping
 - » Updates made to different services will not interfere when they are mapped to different paths
 - `/some/prefix/serviceA/key=(list of URLs)`
 - `/some/prefix/serviceB/key=(list of URLs)`

Server-Side Configuration

```
<beans>
```

```
  <bean id="customTreeCacheMBean"  
    class="org.springframework.jmx.access.MBeanProxyFactoryBean">
```

```
    <property name="objectName">
```

```
      <value>jboss.cache:service=CustomTreeCache</value>
```

```
    </property>
```

```
    <property name="proxyInterface">
```

```
      <value>org.jboss.cache.TreeCacheMBean</value>
```

```
    </property>
```

```
  </bean>
```

```
  <bean id="cacheService"  
    class="app.service.JBossCacheServiceImpl">
```

```
    <property name="cacheMBean"  
      ref="customTreeCacheMBean"/>
```

```
  </bean>
```

```
  <bean id="servicePublisher"  
    class="app.service.ServicePublisher">
```

```
    <property name="cache" ref="cacheService"/>
```

```
  </bean>
```

```
</beans>
```

Client Local Cache Service

- **Step 5: Writing an client-side local JBoss Cache delegate**

```
public class LocalJBossCacheServiceImpl
    implements CacheServiceInterface, ApplicationListener {
    private TreeCache cache;

    public LocalJBossCacheServiceImpl() {
        cache = new TreeCache();
        PropertyConfigurator config =
            new PropertyConfigurator();
        config.configure(cache, "custom-cache-service.xml");
    }

    public void put(String path, Object key, Object val) {
        cache.put(path, key, value);
    }

    public Object get(String path, Object key) {
        return cache.get(path, key);
    }
}
```

Same Cache
configuration as
on server side
(JGroups)

Local Cache Service – Cont'

```
public void onApplicationEvent(ApplicationEvent event) {  
    if (event instanceof ContextRefreshedEvent) {  
        cache.startService();  
    }  
    else if (event instanceof ContextClosedEvent) {  
        cache.stopService();  
    }  
}  
}
```

- ▶ The Client Service holds a reference to a **standalone** copy of JBoss Cache configured from the same configuration file used on the server
 - » JGroups communication layer definitions: UDP etc.
- ▶ All clients and servers access the same Cache!

Client Service Service Definition

```
<bean id="testService"
      class="com.mycompany.AutoDiscoveryProxyFactoryBean">
  <property name="serviceInterface"
    value="com.mycompany.TestServiceInterface" />
  <property name="cache" ref="localCacheService"/>
</bean>
```

- ▶ **AutoDiscoveryProxyFactoryBean:**
 - » Written by app' developer once for all remote services
 - » Extends **HttpInvokerProxyFactoryBean**
 - However, no *serviceUrl* provided
 - » Automatically discovers services
 - Queries the local JBoss Cache
 - Query path is based the bean name
 - » Automatically removes failed services
 - Updates the local JBoss Cache

Client Side Service Discovery

- ▶ **AutoDiscoveryProxyFactoryBean** obtains the list of URLs from our distributed cache:

```
private List getServiceUrls() throws Exception {  
    Set services =  
        (Set)cache.get("/auto/svc/" + beanName, "KEY");  
    if (services == null) {  
        return null;  
    }  
    ArrayList results = new ArrayList(services);  
    Collections.shuffle(results);  
    logger.info("shuffled:" + results);  
    return results;  
}
```

Summary

- ▶ Using the principles described in the last slides, you can deploy Spring Web Services that are:
 - » Load-balanced between the members of your cluster
 - » Stoppable, restartable with dynamic routing
 - » Crash-resistant, failover-enabled



Resources

- ▶ *<http://wiki.jboss.org/wiki/Wiki.jsp?page=JBossSpringIntegration>*
- ▶ *<http://static.springframework.org/spring/docs/1.2.x/reference/jmx.html>*
- ▶ *http://www.javaworld.com/javaworld/jw-10-2005/jw-1031-spring_p.html*
- ▶ *<http://www.jboss.org/products/jbosscache>*





Q&A



Thank You

zvika@tikalk.com