

---

# JBoss Tuning

By : Zvika Markfeld "Tikal"



# Agenda

---

- ▶ JVM
- ▶ Web Container
- ▶ RMI
- ▶ Deployment
- ▶ Session Beans
- ▶ Connection Pools
- ▶ JNDI
- ▶ Transaction Manager
- ▶ Logging

# JVM Tuning – GC Configuration

---

- ▶ We will look at the different garbage collector implementations available in Sun's JDK 1.3+ versions
- ▶ Generational and parallel garbage collection algorithms
- ▶ New in JDK 5.0: Ergonomics

# How Garbage Collectors Work

---

- ▶ *Garbage*: An object that cannot be reached from any pointer in a running program(ref-count=0)
- ▶ Simplistic garbage collection algorithm:
  - » Iterate over all reachable objects, marking them
  - » Any object left over is garbage, reclaim space



# How Garbage Collectors Work

---


- ▶ Problems with simplistic garbage collection:
  - » Inefficient with large memory heaps
  - » Increased GC pauses
  - » Memory Fragmentation
  
- ▶ Hey! Aren't we forgetting something?
  - » Statistically, different objects types have different lifetime:
    - Some objects only live a few milliseconds at a time (local variables, iterators...)
    - Other objects may live days and weeks(data sources, EJB pools)





# Generational Garbage Collector

- ▶ *Generational GC* maintains object pools based on object lifetime
- ▶ Uses different garbage collection algorithms for different pools
  - » Tune your GC based on the application profile
  - » Most applications have high “infant mortality”



# Tuning Generational GC

---

- ▶ The default GC parameters are optimal for client applications
  - » Most server applications are concerned with *Throughput*
    - The total time your JVM spends executing your application and not running the garbage collector
  - » Client applications are concerned with *Pause Times*
    - How long is the noticeable pause the full garbage collection algorithms inflict
- ▶ For best throughput:
  - » Minimize full garbage collections
  - » Keep as many of your object instances in Eden

# Eden

---

- ▶ Young Generation Pool
- ▶ Usually employs a *Minor Garbage Collection*
  - » Uses *Fast Copy Collector*
  - » Copied live objects to a *Survivor Space*
  - » Removes dead objects from Eden
  - » After minor collection, Eden is empty
  - » Quick
  - » Larger memory footprint (copying...)





# Tenured Generation Pools

---

- ▶ Objects are moved from Eden to the *Tenured Generation Pool*:
  - » After surviving several minor collections
  - » When Eden fills up
- ▶ Employ a *Mark-Compact Collector*
  - » Objects are not copied, no need for extra memory allocation
  - » Compacting is much slower than copying
  - » Used for *Major Garbage Collections*
    - When reclaiming short lived objects does not free enough memory

# Configuring Generational Segments

---

- ▶ **-XX:NewRatio=<n>**
  - » Sets ratio between Eden and tenured object pool
  - » Young pool is at maximum  $1/(n+1)$  of total heap
- ▶ **-XX:NewSize, -XX:MaxNewSize**
  - » Sets a fixed size for Eden
  - » Finer granularity



# Configuring Generational Segments

---

- ▶ Note:
  - » The worst case scenario is a young generation pool full of live objects!
  - » When tenured pool is full, a full GC will occur
- ▶ Therefore:
  - » Setting Eden size to more than half the maximum heap size is usually counter productive



# Parallel Garbage Collection

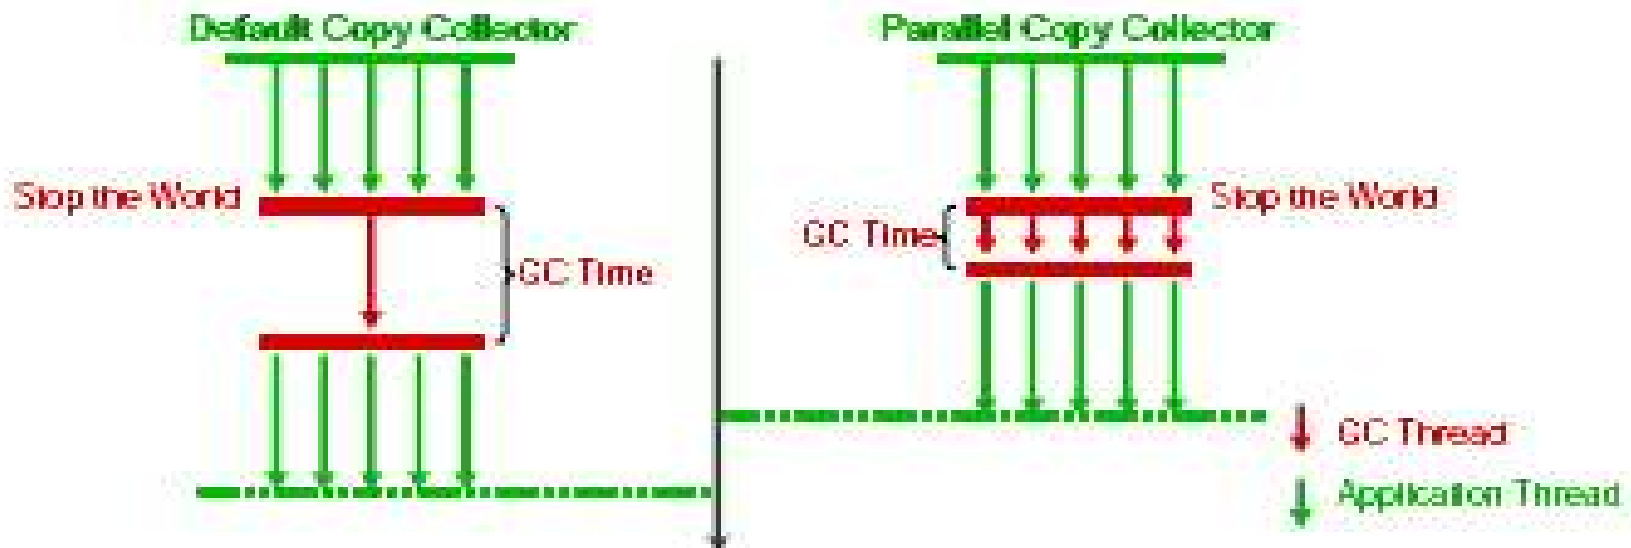
---

- ▶ Most JVMs implement stop-the-world GC algorithms
  - » All running threads are halted
  - » Garbage collection is then activated
- ▶ On Multi-CPU boxes, it makes sense to use all CPUs in sweeping the world clean
- ▶ JDK 1.4.2 introduced two new collectors for that purpose:
  - » *Throughput Collector*
  - » *Low Pause Collector*
- ▶ When running on a multiple CPU box, use JDK 1.4+



# Throughput Collector

- ▶ Parallel Copy Collector
- ▶ Efficient for multi-CPU boxes
- ▶ Employs multiple threads to execute GC
- ▶ Addresses server-side applications



- ▶ `-XX:+UseParallelGC, -XX:ParallelGCThreads=<n>`
  - » Enable the control of the number of GC threads

# Concurrent Low Pause Collector

- ▶ *CMS*: Concurrent mark and sweep collector
- ▶ Collects tenured generations concurrently with application threads
- ▶ Addresses client-side applications



- ▶ **-XX:+UseConcMarkSweepGC**



# Concurrent Low Pause Collector

---

- ▶ By default, the CMS implementation does not compact the tenured pool
- ▶ This might lead to a fragmentation problem, unless using: `-XX:+UseCMSCompactAtFullCollection`.
- ▶ Combine with parallel GC on young generation for maximum throughput
  - » `-XX:+UseParNewGC`

# Other GC Considerations

---

- ▶ Calling `System.gc()` triggers a major collection, negating any benefits from minor collections
  - » Disable with `-XX:+DisableExplicitGC` option
- ▶ RMI DGC subsystem forces a major collection once a minute!
  - » Interval can be controlled with system properties:
    - `-Dsun.rmi.dgc.client.gcInterval=3600000`
    - `-Dsun.rmi.dgc.server.gcInterval=3600000`

# Other GC Considerations

---

- ▶ Default thread stack size on Linux is too large...
  - » Each thread allocates ~8 MB memory
  - » Use: `ulimit -s <nnnn>` to limit stack size
- ▶ Object pooling forces instances to enter the tenured pool, out of reach of the fast copy collector
  - » Negates generational GC
  - » Use Non-Shrinking pools only

# Ergonomics in the JVM 5.0

---

- ▶ Automatic detection of server class machines
  - » More than two CPUs
  - » More than two GB of RAM
- ▶ Upon positive detection:
  - » Throughput garbage collector enabled
  - » Initial heap size of 1 / 64 of physical memory, up to 1GB
  - » maximum heap size of ¼ of physical memory, up to 1GB
  - » Server runtime compiler enabled



# Ergonomics in the JVM 5.0

---

## ▶ *Behaviour Based Tuning*

» `-XX MaxGCPauseMillis=<nnn>`

Maximum pause time goal: a hint to the GC that pause times of <nnn> milliseconds or less are desired

» `-XX:GCTimeRatio=<nnn>`

Application throughput goal: the ration between the time spent on GC and the time spent on application

▶ The JVM guesses the rest of the parameters...



# Web Containers – Connectors

- ▶ JBossWeb container can be found under:  
`deploy/jbossweb-tomcatXXX.sar`
- ▶ The connection configuration is found in `server.xml`
- ▶ *Connectors* are thread pools that accept inbound requests and process them

```
<!-- A HTTP/1.1 Connector on port 8080 -->  
<Connector
```

```
    port="8080"  
    address="${jboss.bind.address}"  
    maxThreads="250"  
    minSpareThreads="25"  
    maxHttpHeaderSize="8192"  
    emptySessionPath="true"  
    enableLookups="false"  
    redirectPort="8443"  
    acceptCount="100"  
    connectionTimeout="20000"  
    disableUploadTimeout="true"/>
```



# Web Containers – Connectors

---

- ▶ **AcceptCount:**  
Length of the incoming request queue(when there are no processors available)
- ▶ **ConnectionTimeout:**  
How long to wait before a URI is received from the stream (the default is 20 seconds)
  - » This avoid problems where a client opens a connection and does not send any data
- ▶ **MaxThreads :**  
The maximum number of threads in the pool, this is a strict pool.
  - » Rule of thumb: set to 25% more than your maximum expected load (concurrent hits coming in at once)

# Web Containers – Connectors

---

- ▶ **EnableLookups :**  
Whether to perform a reverse DNS lookups to prevent snoofing
  - » Might cause problems when a DNS is ‘misbehaving’
  - » Turn off when you implicitly trust all clients
- ▶ **MinSpareThreads :**  
“On start up, always keep at least this many threads waiting idle”
  - » Rule of thumb: Set to a little more than your normal load
- ▶ **MaxSpareThreads :**  
“If we ever go above minSpareThreads, always keep this number of threads waiting idle”
  - » Rule of thumb: Set to a little more than your peak load

# Web Containers – Valves

- ▶ Remove any unnecessary valves and logging
  - » For example, if not using JBoss security, remove the security valve

```
<!--  
  <Valve className="org.apache.catalina.valves.RequestDumperValve" />  
  
  <Valve  
    className="org.apache.catalina.valves.FastCommonAccessLogValve"  
      prefix="localhost_access_log." suffix=".log"  
      pattern="common"  
    directory="${jboss.server.home.dir}/log"  
      resolveHosts="false" />  
-->
```

# Web Containers – JSP Optimization

- ▶ Jasper configuration found in: `conf/web.xml`
- ▶ Configured by default to `development` mode
- ▶ This means the jsp page might be checked for modification on *every* access

```
<servlet>
  <servlet-name>jsp</servlet-name>
  <servlet-class>org.apache.jasper.servlet.JspServlet</servlet-
class>
  <init-param>
    <param-name>fork</param-name>
    <param-value>>false</param-value>
  </init-param>
  <init-param>
    <param-name>development</param-name>
    <param-value>>false</param-value>
  </init-param>
  <init-param>
    <param-name>checkInterval</param-name>
    <param-value>60</param-value>
  </init-param>
  <load-on-startup>3</load-on-startup>
</servlet>
```

# Web Containers – JSP Optimization

---

- ▶ **Development=false:**  
Turns off checking for changes in JSPs
- ▶ **CheckInterval:**  
While in dev' mode, sets JSP check interval
- ▶ JSPs can also be precompiled to servlets to avoid delay on first access

# Remote Invocations

- ▶ By default, JBoss creates a new thread for every remote invocation(RMI) request
  - » Not efficient on a large systems
  - » Inadvisable in the case of performance or traffic peaks
- ▶ Instead, use pooled invokers:
  - » Edit `conf/standardjboss.xml`:
  - » Change all of the proxy bindings to the pooled invoker by changing every XML fragment reading:

```
<!--  
<invoker-  
mbean>jboss:service=invoker,type=jrmp</invoker-mbean>  
-->
```

```
<invoker-  
mbean>jboss:service=invoker,type=pooled</invoker-mbean>
```



# Deployment Scanner

- ▶ The deployment scanner scans for new deployments every 5 seconds
- ▶ This eats up cycles, especially on systems with a slow file system
- ▶ Edit `conf/jboss-service.xml`:

```
<!-- An mbean for hot deployment/undeployment of archives.
-->
```

```
<mbean
code="org.jboss.deployment.scanner.URLDeploymentScanner"
name="jboss.deployment:type=DeploymentScanner,flavor=URL">
...
```

```
<attribute name="ScanPeriod">5000</attribute>
...
</mbean>
```

# Stateless Session Beans

- ▶ EJB stateless session beans operate according to a spec-dictated pooling model
- ▶ If you find that you need more than the default (10) instances, set minimum pool size:

» Edit `conf/standardjboss.xml`:

```
<container-pool-conf>  
  <MaximumSize>100</MaximumSize>  
</container-pool-conf>
```



```
<container-pool-conf>  
  <MinimumSize>100</MinimumSize>  
  <MaximumSize>100</MaximumSize>  
  <strictMaximumSize/>  
  <strictTimeout>30000</strictTimeout>  
</container-pool-conf>
```

- ▶ Usually, we wouldn't want these pools growing and shrinking due to memory-fragmentation
- ▶ From a performance standpoint, the number should be big enough to serve all your requests with no blocking

# Datasource Configuration

---

- ▶ Use XA connections only when necessary
  - » XA connections have performance issues
- ▶ Use database specific "ping" support where available for "check-connection"
- ▶ Use database-specific driver fail-over support rather than checking connections at all
  - » Remember that not all tuning options may be available in your environment, we're talking optimal here...



# JNDI

- ▶ It is a common mistake to configure a provider url on the server, as in:

```
java.naming.provider.url=jnp://localhost:1099
```

- ▶ This forces JNDI to use local sockets for local access
- ▶ Inside the server, jndi should be done using internal method calls
  - » This occurs when there is no provider url
  - » On the client side, however, this forces a multicast discovery of HAJNDI
  - » You may disable this behaviour by setting:  
`jnp.disableDiscovery=true`

# Transaction Manager

- ▶ The transaction manager configuration can be found in: `conf/jboss-service.xml`
- ▶ The transaction timeout helps breaking blocked threads

```
<mbean code="org.jboss.tm.TransactionManagerService"
      name="jboss:service=TransactionManager"
      xmbean-dd="resource:xmdesc/TransactionManagerService-xmbean.xml">
  <attribute name="TransactionTimeout">300</attribute>
  ...
</mbean>
```

- ▶ **TransactionTimeout:**  
How long a transaction lasts before the thread gets interrupted and the transaction marked for rollback
  - » The default is 5 minutes


# Log4J

- ▶ Logging has a profound effect on performance
  - » Setting logging level to TRACE can bring JBoss to a crawl
  - » Changing it to ERROR can improve speed dramatically
- ▶ By default, JBoss logs INFO-level messages to the console and to `server.log`
  - » Consider not logging to the console
  - » Consider changing the log level to ERROR / WARN
  - » Add a category filter for your Java class hierarchy
- ▶ NOTE:
  - » JBoss watches its Log4J configuration file for changes
  - » You can always change configuration at runtime



# Log4j

```
<root>  
  <appender-ref ref=CONSOLE"/>  
  <appender-ref ref="FILE"/>  
</root>
```



```
<root>  
  <priority value="ERROR" />  
  <appender-ref ref="FILE"/>  
</root>
```



# Q&A



***Thank you***

***zvika@tikalk.com***